

# Network QoS

## 371-2-0213

### Lecture 4

Gabriel Scalosub

# Outline


- 1 Recap
  - DiffServ Fundamentals
  - Competitive Analysis
- 2 Competitive Buffer Management with Commitments
  - Model and Preliminary Observations
  - A Competitive Algorithm
    - Upper Bound Preliminaries
    - Algorithm ON
    - Competitive Analysis of ON
  - Simulation and Conclusions

# DiffServ - Key Components


- Main motivation: IntServ handicaps
  - per-flow state
  - complex reservation mechanisms
  - requires large-scale deployment
- Main concern: being “better than best-effort”
- Main design concept: locality of decision
- Key components:
  - middle ground between IntServ and best-effort
  - divide traffic into small number of *forwarding classes*
    - E.g., Gold, Silver, Bronze, Best-Effort
    - class encoded in packet IP header
    - resources allocated per-class (aggregate)
  - based on Service Level Agreements (SLAs)
  - mechanisms
    - marking
    - scheduling (priority queuing, WFQ)
    - buffer management (AQM/RED)

# SLAs and Marking

- Specification of expected traffic per-class
  - token-bucket envelope
    - Committed/Peak Information Rate (CIR/PIR)
    - Committed/Peak Burst Size (CBS/PBS)
- Three color marking
  - Done using a dual token bucket
  - E.g., two-rate three color marking



# DiffServ Codepoint (DSCP)



- Default PHB:
  - DSCP:  $\langle 000000 \rangle$  (best-effort)
- Assured Forwarding (AF) PHBs:
  - DSCP:  $\langle 001xx0 \rangle$ ,  $\langle 010xx0 \rangle$ ,  $\langle 011xx0 \rangle$ ,  $\langle 100xx0 \rangle$
  - at least two different forwarding classes
  - per class bandwidth allocation
    - implemented using scheduling (e.g., WFQ)
    - each link can determine its own allocation
  - 3 drop priorities per class (e.g., via RED)
- Expedited Forwarding (EF) PHB:
  - DSCP:  $\langle 101110 \rangle$
  - high-priority queue (10-30% of link capacity)

# Competitive Analysis - Definition

- Given an instance  $I$  of an optimization problem  $\mathcal{P}$ , denote by  $\text{OPT}(I)$  the value of an optimal feasible solution for  $I$ .
- In the online setting, the input to the problem is made available in parts.
- An online algorithm  $A$  is said to be  $c$ -competitive for problem  $\mathcal{P}$  if for every instance  $I$  of  $\mathcal{P}$ ,  $A(I)$  satisfies:
  - $A(I) \leq c \cdot \text{OPT}(I) + \alpha$  (minimization problem  $\mathcal{P}$ )
  - $A(I) \geq \frac{1}{c} \cdot \text{OPT}(I) - \alpha$  (maximization problem  $\mathcal{P}$ )where  $\alpha \geq 0$  is some additive term independent of  $I$ .
- Common to assume that  $I$  is generated by an *adversary*
- The online problem is then viewed as a *game*:
  - Adversary (produces  $I$  and an optimal solution to  $I$ ), vs.
  - Algorithm

# Model

- Provider's viewpoint: fulfill its end of the SLA
- Focus on a single AF PHB class
- Traffic model
  - committed traffic (green):  $(r, B)$  token envelope
  - excess traffic (yellow): arbitrary
  - interleaved
- Queue model
  - single FIFO queue  $Q$ 
    - buffer of size  $B_Q \geq B$
    - service rate  $r_Q \geq r$
    - preemptive: may drop enqueued packets

## Feasibility

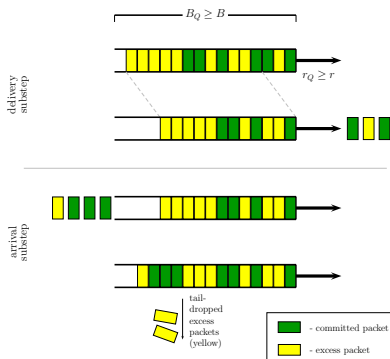
Never drop committed packets

## Goal

Maximize number of excess packets forwarded

# Time Model and Simplifying Assumptions

- Slotted time
  - delivery/forwarding substep:  $\leq r_Q$  forwarded
  - arrival substep: packets may be dropped/accepted



- Simplifying assumption:
  - Uniform size packets (WLOG, unit size)
  - Unit rates (i.e.,  $r_Q = r = 1$ )



# Resource Augmentation is Required

## Theorem

*Any online algorithm ALG using  $B_Q \leq B$  cannot be competitive*

## Proof.

- Assume  $B_Q = B$ .
- $t = 0$  arrival:  $\boxed{Y}, \boxed{G}$
- $t = 1$  forwarding: two cases
  - $\boxed{Y}$  dropped: ALG cannot be competitive
    - since no more yellow packets may arrive
  - $\boxed{Y}$  forwarded, and  $\boxed{G}$  enqueued
    - $t = 2$  arrival: burst of  $B$   $\boxed{G}$ s
    - ALG has to drop  $\boxed{G}$

**Infeasible!!**



# Preliminaries and Basic Concepts

- Use a buffer of size  $(1 + \varepsilon)B$ 
  - OPT will use a buffer of size  $B$
- Notation
  - $B_A(t)$ : set of packets in the buffer at time  $t$  under algorithm  $A$
  - $d_t^A(p)$ : buffer position of packet  $p$  at time  $t$  under algorithm  $A$
- Lower bound on OPT buffer occupancy
  - use a *simulator* SIM
  - SIM has the same resources as OPT (rate, buffer)
  - SIM ignores all yellow packets
  - properties of SIM:
    - $\forall t$  and  $\forall$  green  $p$  that arrived by  $t$ ,  $d_t^{\text{SIM}}(p) \leq d_t^{\text{OPT}}(p)$
    - $\forall t$   $|B_{\text{SIM}}(t)| = k$  implies  $|B_{\text{OPT}}(t)| \geq k$
- A naïve approach:
  - maintain two queues
    - green (size  $B$ ) and yellow (size  $\varepsilon B$ )
  - always give priority to green queue
    - green queue is equivalent to SIM
  - but...

# Concept of Lag

- Lag of a green packet at time  $t$  under algorithm  $A$ :

$$\text{lag}_t^A(p) = d_t^A(p) - d_t^{\text{SIM}}(p)$$

- for green  $p, p' \in B_A(t)$ , lag is monotone non-dec. in arrival
- Context: after forwarding substep or arrival substep?
  - Depends

## $\delta$ -lag property

Algorithm  $A$  satisfies the  $\delta$ -lag property if  $\forall t, p, \text{lag}_t^A(p) \leq \delta$

## Lag of an algorithm

The lag of algorithm  $A$  at time  $t$  is  $\phi_A(t) = \max_{p \in B_A(t)} \text{lag}_t^A(p)$

- $\phi_A(t)$  determined by last green packet in  $B_A(t)$ 
  - FIFO + monotonicity of lag
- For now: consider lag at end of forwarding substep

# Algorithm ON

---

**Algorithm 1** ON: upon the arrival of a new packet  $p$

---

- 1: **if**  $p$  is yellow **then**
  - 2:     accept if there's room
  - 3: **else** ▷  $p$  is green
  - 4:     Drop as few yellow packets from the tail of the queue such  
      that the new packet will have lag at most  $\varepsilon B$
  - 5:     Accept  $p$
  - 6: **end if**
- 

- omit ON sub/super-scripts: use  $\text{lag}_t(p)$ ,  $\phi(t)$

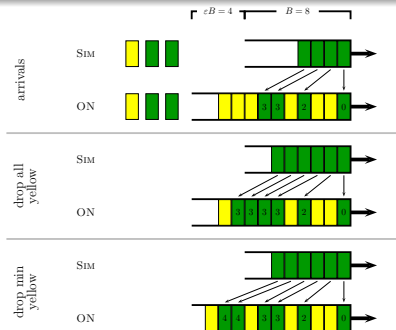
# Feasibility of ON

## Theorem (Feasibility and lag)

At any time  $t$  algorithm ON

- accepts all green packets
- always holds  $\leq (1 + \varepsilon)B$  packets
- satisfies the  $\varepsilon B$ -lag property

“Proof by picture”  
(formally, by induction)



# Competitive Ratio of ON

## Theorem (Competitive ratio)


*Algorithm ON is  $\frac{\varepsilon}{1+\varepsilon}$ -competitive*

- Analysis in a nutshell
  - identify “reset” events
  - overflow (yellow dropped) occurs only between resets
    - overflow intervals
  - at least  $\varepsilon B$  yellow are “safe” since last reset
  - many green accepted by SIM
    - OPT must deal with them too
    - has little room for many yellow

# Notation and Definitions

- From now on: consider lag at end of delivery substep
- Some additional notation and definitions
  - $A_t^G$ : the set of green packets arriving at  $t$
  - $A_t^Y$ : the set of yellow packets arriving at  $t$
  - for simplicity: assume each such set is handled as a batch
    - first green, then yellow
- Safe packets
  - Yellow  $p \in B_{\text{ON}}(t)$  *turns safe at  $t + 1$* 
    - $t$  is minimal s.t.  $A_t^G \neq \emptyset$ , and  $p$  is not dropped at  $t$
  - $S_t$ : set of packets turning safe at  $t$
  - for every time interval  $I$ ,  $S(I) = \cup_{t \in I} S_t$
- Reset
  - ON is *reset at  $t$*  if  $\phi(t) = 0$ 
    - queue idle at  $t$  implies reset at  $t$
    - no green in buffer at  $t$  implies reset at  $t$

# Overload Intervals: Following the Lag Process



- Between overload intervals, ON does at least as good as OPT
- Suffices to analyze performance in overload intervals
- Any two overload intervals are independent
- Focus on a single overload interval,  $I$



# Understanding Changes in Lag

## Lemma

For any non-reset  $t \in I$  s.t.  $\phi(t) > 0$ ,  

$$\phi(t) = \phi(t-1) + |S_t| - (1 - \mathbb{1}_{\text{SIM}}(t))$$

## Proof.

- Case 1:  $A_{t-1}^G = \emptyset$ 
    - necessarily  $S_t = \emptyset$  (by definition of turning safe)
    - $\phi(t) > 0$ :
      - last green  $p \in B_{\text{ON}}(t-1)$  still in buffer at the end of  $t$
      - $p$  advances
        - $\mathbb{1}_{\text{SIM}}(t)$  places in SIM
        - one place in ON
      - $\text{lag}_t(p) = \text{lag}_{t-1}(p) - (1 - \mathbb{1}_{\text{SIM}})$
- $\Rightarrow \phi(t) = \phi(t-1) - (1 - \mathbb{1}_{\text{SIM}}(t)) = \phi(t-1) + |S_t| - (1 - \mathbb{1}_{\text{SIM}}(t))$

# Understanding Changes in Lag

## Lemma

For any non-reset  $t \in I$  s.t.  $\phi(t) > 0$ ,  

$$\phi(t) = \phi(t-1) + |S_t| - (1 - \mathbb{1}_{\text{SIM}}(t))$$

## Proof.

- Case 2:  $A_{t-1}^G \neq \emptyset$ 
  - consider
    - last green  $p \in B_{\text{ON}}(t-1)$  (end of  $t-1$ ), and
    - last green  $p' \in B_{\text{ON}}(t-1)$  (end of delivery substep of  $t-1$ )
    - $p'$  exists since otherwise reset at  $t-1$ , and hence in  $t$
  - relative position of  $p$  and  $p'$  at end of  $t-1$ 
    - $p$  is exactly  $|A_{t-1}^G| + |S_t|$  positions behind  $p'$  in ON
    - $p$  is exactly  $|A_{t-1}^G|$  positions behind  $p'$  in SIM
  - $\text{lag}_t(p) = \text{lag}_{t-1}(p') + |S_t| - (1 - \mathbb{1}_{\text{SIM}}(t))$ 
    - $p$  not sent at  $t$  since  $\phi(t) > 0$

$$\Rightarrow \phi(t) = \phi(t-1) + |S_t| - (1 - \mathbb{1}_{\text{SIM}}(t))$$



# Understanding Load on OPT During Overload

## Lemma

*SIM delivers  $|I| - |S(I)|$  green packets during  $I$*

## Proof.

- $R_I$ : set of green packets delivered by SIM during  $I$
- $\Delta(t) = \phi(t) - \phi(t-1) = |S_t| - (1 - \mathbb{1}_{\text{SIM}}(t))$
- On one hand

$$\begin{aligned} \sum_{t \in I} \Delta(t) &= \sum_{t \in I} [|S_t| - (1 - \mathbb{1}_{\text{SIM}}(t))] \\ &= \sum_{t \in I} |S_t| - \sum_{t \in I} 1 + \sum_{t \in I} \mathbb{1}_{\text{SIM}}(t) \\ &= |S(I)| - |I| + |R_I| \end{aligned}$$

- On the other hand,  $\sum_{t \in I} \Delta(t) = 0$ 
  - telescopic sum

$$\Rightarrow |R_I| = |I| - |S(I)|$$



# Proof of Competitive Ratio

## Theorem (Competitive ratio)

*Algorithm ON is  $\frac{\varepsilon}{1+\varepsilon}$ -competitive*

Proof (focus on single overload interval  $I$ ).

- Recall: OPT has rate  $r = 1$  and a buffer of size  $B$ 
  - OPT deals with at most  $r|I| + B = |I| + B$  packets altogether
  - has to deal with at least  $|R_I|$  green handled by SIM

$\Rightarrow$  # yellow packets handled by OPT during  $I$  is at most  
 $(|I| + B) - |R_I| = (|I| + B) - (|I| - |S(I)|) = |S(I)| + B$

- $|S(I)| \geq \varepsilon B$ 
  - definition of overflow interval
  - $\phi(t)$  unit-increase implies a yellow packet turning safe
- Competitive ratio is at least

$$\frac{|S(I)|}{|S(I)| + B} \geq \frac{\varepsilon B}{\varepsilon B + B} = \frac{\varepsilon}{1 + \varepsilon}$$

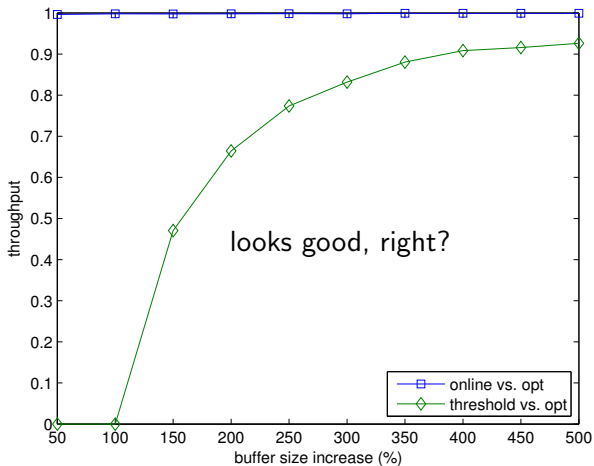


# Simulation Results

- Traffic
  - bursty MMPP traffic
    - (two) color marked using dual token bucket
  - best-effort Poisson traffic (cross-traffic)
    - zero-rate committed
  - interleaved
- Contending protocols
  - Threshold
    - accept yellow only if buffer occupancy is below threshold  $T$
    - feasibility implies  $T = \varepsilon B$  (should be  $\leq$ )
    - commonly used policy (single-step RED)
  - Naïve protocol
    - 2-queues priority queuing
    - not FIFO, but
    - upper bound on  $\text{OPT}$
    - serves for normalization


# Simulation Results

- Single MMPP aggregate, yellow  $\sim 30\%$




# Simulation Results

- MMPP+Poisson, yellow also during OFF, yellow  $\sim 40\%$



# Simulation Results

- MMPP+Poisson, yellow also during OFF, yellow  $\sim 50\%$





# Open Questions and Extensions

- Lower bounds
  - for  $\varepsilon \leq 1$ 
    - still with resource augmentation
    - CR is at least  $\varepsilon$
- What's the right answer
  - lower bounds for  $\varepsilon > 1$ ?
  - if we double the buffer ( $\varepsilon = 1$ ), do we obtain OPT?
  - closing gap for  $\varepsilon \in (0, 1]$
- Models that allow dropping “some” committed packets
- Network perspective
  - end-to-end throughput
  - simple topologies (line, tree, ...)

- 
- References
    - Patt-Shamir, Scalosub and Shavitt, *Competitive Analysis of Buffer Policies with SLA Commitments*, ICNP 2008

# A Couple of Lessons Learnt

- FIFO
  - might delay the delivery of “important” packets
  - might have a delayed effect
  - one has to be careful with that
- Competitive approach
  - main focus: find an upper bound on  $OPT!!!$
  - in this analysis: an averaging argument
    - long enough intervals
    - compare  $ALG$  and  $OPT$  over such intervals
  - later in the course: we’ll see other techniques
- Simulation results may be misleading
  - must understand the traffic
  - if it’s too good to be true, there’s probably a bug...
  - must look at the logs to verify soundness
  - simulations do not “prove” anything
    - traffic might be too “easy” / “hard”
    - difficult to represent real life
    - usually serves to further validate analytic results