

# Network QoS

## 371-2-0213

### Lecture 6

Gabriel Scalosub

# Outline

- 1 Competitive Weight-based Buffer Management
  - Basic Model
  - Upper Bound
  - Lower Bound
  
- 2 Going Beyond: Algorithms and Models
  - Better Algorithms
  - Other Models

# Model

- Commitments model we saw: managing a single AF PHB
- Now: a more general DiffServ Model
  - Managing multiple AF PHBs / EF classes
  - Differentiate between packets corresponding to different classes
  - E.g.: EF ">" AF41 ">" AF43
    - scheduling: link access (delay / rate)
    - buffer management: drop rate
- System model: slotted time
- Traffic model
  - uniform size packets (WLOG, size is 1)
  - for each packet  $p$  we have
    - $a(p)$ : arrival time
    - $w(p)$ : packet *weight*
    - $s_A(p)$ : forwarding/delivery time by algorithm  $A$
- Packet weight gives means for differentiation
  - can be considered as part of PHB

# Model

- Queue model
  - Single FIFO queue  $Q$ 
    - buffer of size  $B$
    - service rate  $r$
    - preemptive: may drop enqueued packets

## Goal

Maximize overall weight of forwarded packets

- What algorithm would you pick?
- Assumptions and notation:
  - $B_A(t)$ : set of packets in the buffer of algorithm  $A$  at time  $t$
  - $S_A(t)$ : set of packets sent by algorithm  $A$  at time  $t$
  - $w(P)$ : overall weight of set of packets  $P$ 
    - i.e.,  $w(P) = \sum_{p \in P} w(p)$
  - for simplicity, assume  $r = 1$

# The Greedy Algorithm

---

**Algorithm 1** G: upon the arrival of  $p$  at time  $t$

---

```
1: if  $|\mathcal{B}_G(t)| < B$  then
2:   accept  $p$ 
3: else
4:    $q \leftarrow \arg \min \{w(q') \mid q' \in \mathcal{B}_G(t)\}$ 
5:   if  $w(p) \leq w(q)$  then
6:     drop  $p$ 
7:   else
8:     drop  $q$ 
9:     accept  $p$ 
10:  end if
11: end if
```

---

- Notation:

- $A(t)$ : packets arriving at  $t$
- $D_G(t)$ : packets dropped by G at  $t$

## Competitive Analysis of Greedy - Upper Bound

- $V(t)$ :  $B$  highest weight packets arriving during  $[t + 1, t + B]$

### Lemma

*For any time  $t$ ,*

$$w(\mathcal{B}_G(t)) + w(V(t)) \leq \sum_{i=1}^B w(s_G(t + i)) + w(\mathcal{B}_G(t + B))$$

- LHS: current buffer at  $t$  + best potential arrivals by  $(t + B)$
- RHS: actual deliveries by  $(t + B)$  + future buffer at  $(t + B)$

# Competitive Analysis of Greedy - Upper Bound

- $V(t)$ :  $B$  highest weight packets arriving during  $[t + 1, t + B]$

## Lemma

For any time  $t$ ,

$$w(\mathcal{B}_G(t)) + w(V(t)) \leq \sum_{i=1}^B w(s_G(t + i)) + w(\mathcal{B}_G(t + B))$$

## Proof.

- $Y_0^t$ :  $\mathcal{B}_G(t) \cup V(t)$
- $Y_i^t = Y_{i-1}^t \cup A(t + i) \setminus D_G(t + i)$  (for  $i \geq 1$ )
  - By definition of  $G$ ,  $w(Y_{i-1}^t) \leq w(Y_i^t)$   
 $\Rightarrow w(Y_0^t) \leq w(Y_B^t)$
  - $Y_B^t \subseteq \mathcal{B}_G(t) \cup \bigcup_{i=1}^B A(t + i)$
  - none of the packets in  $Y_B^t$  are discarded by  $t + B$
  - packets in  $Y_B^t$  either sent by  $t + B$ , or in  $\mathcal{B}_G(t + B)$



# Competitive Analysis of Greedy - Upper Bound

## Theorem

*Algorithm G is 2-competitive*

## Proof.

- Partition time into intervals of length  $B$ 
  - $I_k = [kB, (k+1)B)$ ,  $k = 0, 1, \dots$
- $U(k)$ : set of packets accepted by  $\text{OPT}$  during  $I_k$
- $|I_k| = B$ 
  - $|U(k)| \leq 2B$
  - $w(U(k)) \leq 2w(V(kB))$



# Competitive Analysis of Greedy - Upper Bound

## Theorem

*Algorithm G is 2-competitive*

## Proof.

- Previous lemma:

$$w(V(kB)) \leq \sum_{i=1}^B w(s_G(kB+i)) + w(\mathcal{B}_G(kB+B)) - w(\mathcal{B}_G(kB))$$

- for  $T = \max \{a(p)\} + B$


$$\sum_k w(V(kB)) \leq \sum_t w(s_G(t)) + \underbrace{w(\mathcal{B}_G(T))}_0 - \underbrace{w(\mathcal{B}_G(0))}_0$$

$$\Rightarrow w(OPT) \leq 2 \sum_k w(V(kB)) \leq 2w(G)$$



# Competitive Analysis of Greedy

- The analysis of  $G$  is essentially tight
  - Consider the traffic:



- $w(G) = B(1 + \alpha)$
  - $w(OPT) = 2B\alpha$
  - the ratio is  $\frac{2B\alpha}{B(1+\alpha)} = \frac{2\alpha}{1+\alpha} = 2 - \frac{2}{1+\alpha}$
  - tends to 2 as  $\alpha$  goes to infinity
- Uses only 2 distinct values


# Competitive Analysis - Deterministic Lower Bound

## Theorem

*No deterministic algorithm has competitive ratio better than 1.281*

## Proof.

- Fix some deterministic algorithm ALG
- Consider traffic similar to previous
  - continues until earliest time  $t \leq B$ :
    - ALG sends an  $\alpha$ -packet at  $(t+1)$ ,  
or
    - $t = B$
- Scenario one:
  - sequence ends
    - ALG gains  $t + t\alpha$
    - OPT would have gained  $B + t\alpha$



# Competitive Analysis - Deterministic Lower Bound


## Theorem

*No deterministic algorithm has competitive ratio better than 1.281*

## Proof.

- Scenario two:
  - sequence continues similar to greedy's tight example
    - ALG gains  $t + B\alpha$
    - OPT would have gained  $(t + B)\alpha$
- Competitive ratio of ALG is at least

$$\max \left\{ \frac{B + t\alpha}{t + t\alpha}, \frac{(t + B)\alpha}{t + B\alpha} \right\}$$



## Competitive Analysis - Deterministic Lower Bound

### Theorem

*No deterministic algorithm has competitive ratio better than 1.281*

### Proof.

- The minimum (over  $t$ ) is obtained when equality holds
  - Solving (for  $t$ ):  $t_0(B, \alpha) = B \frac{\sqrt{(\alpha-1)^2 + 4\alpha^3} - \alpha + 1}{2\alpha^2}$
  - define  $t_0^*(\alpha) = t_0(B, \alpha)/B$
  - competitive ratio is at least  $\frac{1 + \alpha t_0^*(\alpha)}{(1 + \alpha) t_0^*(\alpha)}$
- solving (numerically, for  $\alpha$ ), this is maximized for  $\alpha \approx 4.01545$
- Competitive ratio of ALG cannot be better than 1.281



# Better Algorithms: What Would You Do?

- Recall Greedy:
- Improvement options:
  - consider *complex  $p$ - $q$  relationships*
    - drop  $q$  only if  $w(q) \ll w(p)$
    - not just comparison-based – actual values matter!!
  - consider the overall *present “state”* of the buffer
    - how much value is currently in the buffer?
    - do we currently have enough to “cover” for not preempting?
    - proactive dropping? (avoid delaying others...)
  - consider the present state and *history*
    - how much value did we gain recently?
    - did we gain enough to “cover” for not preempting?
  - admission control*
    - drop even if the buffer isn't full? (proactive...)
    - may depend on / be combined with all of the above
  - toss coins...*

---

**Algorithm 1** G: upon the arrival of  $p$  at time  $t$ 

---

```
1: if  $|\mathcal{B}_G(t)| < B$  then
2:   accept  $p$ 
3: else
4:    $q \leftarrow \arg \min \{w(q') \mid q' \in \mathcal{B}_G(t)\}$ 
5:   if  $w(p) \leq w(q)$  then
6:     drop  $p$ 
7:   else
8:     drop  $q$ 
9:     accept  $p$ 
10:  end if
11: end if
```

---

# How Was This Useful?

- Detailed results we saw are due to:
  - upper bound: Mansour, Patt-Shamir, Lapid (2000)
  - lower bound: Kesselman, Lotker, Mansour, Patt-Shamir, Schieber, Sviridenko (2001)
- Many papers tried to close the gap
  - many algorithms!!
- Most of the focus: 2-valued case
  - matching upper bound finally by Englert&Westermann (2008)
- What about randomized algorithms?
  - strictly better than deterministic...

deterministic		randomized	
upper	lower	upper	lower
	1.282		
$2 - \frac{2}{\alpha+1}$			
1.894			
1.544			
1.304			
		1.25	1.197
1.282			

## Non-preemptive, 2-valued FIFO

- Actually, the first model considering packet weights
  - Aiello, Mansour, Rajagopalan, Rosén (2000)
- Focus solely on admission control
- greedy is bad: tight  $1/\alpha$ -competitive
  - upper bound: every 1-packet mapped to an  $\alpha$ -packet
  - lower bound: non-preemption implies algorithm may only accept 1-packets
- lower bound:  $\frac{\alpha}{2\alpha-1}$
- matching upper bound by Andelman, Mansour, Zhu (2003)



# General Values

- Many more papers...

deterministic		randomized	
upper	lower	upper	lower
4			
2			
	1.282		
	1.414		
	1.419		
	1.434		
1.983			
1.75			
		1.75	1.25
1.732			

## More Models

- Loss-bounded model
  - measure: packets dropped (instead of delivered)
  - Kesselman&Mansour (2003)
  - constant competitive
    - also implies good competitive ratio in “standard” model
- Multiple Queues
  - multiple output-queues with shared memory
    - unit-values (not really “DiffServ”)
    - goal is doing load-balancing (e.g., longest-queue-drop)
    - constant competitive
  - multiple input-queues
    - essentially a scheduling problem: which IQ to schedule
    - goal is doing load-balancing (e.g., longest-queue-first)
    - mostly constant competitive
- Many more...