

Network QoS

371-2-0213

Lecture 8

Gabriel Scalosub

Outline

- 1 Classical Scheduling Algorithms
 - Introduction and Makespan
 - Interval Scheduling
 - Aftermath

- 2 Final Project: Some Thoughts

Scheduling Preliminaries

- Over 50 years of research
- Model:
 - n jobs J_j , $j = 1, \dots, n$
 - m machines M_i , $i = 1, \dots, m$
 - p_{ij} : job J_j 's processing time on machine M_i
 - jobs should be scheduled on machines
 - various constraints/assumptions
 - objective function is optimized
- Examples:
 - schedule jobs on m machines
 - variable-length jobs (same on all machines)
 - goal: minimize maximum completion time (Makespan)
 - single machine: trivial...
 - two machines: NP-hard (subset-sum)

Machine Scheduling Notation: $\alpha \mid \beta \mid \gamma$

- Machine environment α :
 - m : number of machines
 - P : parallel identical machines
 - Q : parallel machines with speeds, $p_{ij} = p_j/s_i$
 - job processing requirement p_j , machine speed s_i (independent)
 - R : parallel unrelated machines
 - ...
- Job characteristics β :
 - r_j : release time
 - d_j : deadline
 - pmtn: preemption is allowed
 - a job assigned to a machine can be stopped, and rescheduled later
 - prec: precedence constraints
 - specified by an underlying DAG on the jobs
 - ...

Machine Scheduling Notation: $\alpha \mid \beta \mid \gamma$ (Cont.)

- Objective function γ :
 - C_{\max} : minimize maximum completion time of a job
 - job J_j 's completion time C_j
 - L_{\max} : minimize maximum lateness of a job
 - job J_j 's lateness $L_j = C_j - d_j$
 - $\sum_j U_j$: minimize number of late jobs
 - job J_j 's lateness indicator $U_j = \mathbb{1}_{C_j > d_j}$
 - $\sum_j w_j \bar{U}_j$: maximize weight of jobs scheduled by their deadlines
 - job J_j 's non-lateness indicator $\bar{U}_j = 1 - U_j$
 - ...
- Examples:
 - $1 \mid d_j, r_j \mid \sum_j w_j \bar{U}_j$
 - one machine ("queue"), maximize weight of jobs scheduled by their deadlines, with release times.
 - similar to the problem we saw last week...
 - $Q2 \mid \text{pmtn, prec, } d_j \mid \sum_j L_j$
 - 2 machines (different speeds), preemption allowed, precedence constraints, minimize sum (equiv. – average) of latenesses

Machine Scheduling Notation: $\alpha \mid \beta \mid \gamma$ (Cont.)

- One more example:
 - $Pm \parallel C_{\max}$
 - m identical machines, minimize the *makespan*
 - coming up next...

Examples of Problems, or “The Scheduling Zoo”

<http://www.informatik.uni-osnabrueck.de/knust/class/>,

<http://schedulingzoo.lip6.fr/>

Parallel machine problems without preemption

• maximal polynomially solvable:

$P p_i = p; outtree; r_i C_{max}$	Brucker et al. (1977)
$P p_i = p; tree C_{max}$	Hu (1961), David & Linton (1976)
$Q p_i = p; r_i C_{max}$	Assignment-problem
$Q2 p_i = p; chains C_{max}$	Brucker et al. (1999)
$P p_i = 1; chains; r_i L_{max}$	Deur et al. (1998), Baptiste et al. (2004)
$P p_i = p; intree L_{max}$	Brucker et al. (1977), Monma (1982)
$P2 p_i = p; prec L_{max}$	Garey & Johnson (1976)
$P2 p_i = 1; prec; r_i L_{max}$	Garey & Johnson (1977)
$P p_i = 1; outtree; r_i \sum C_i$	Brucker et al. (2002), Huo & Leung (2005)
$P p_i = p; outtree \sum C_i$	Hu (1961)
$P2 p_i = 1; prec; r_i \sum C_i$	Baptiste & Tiskovsky (2004)
$P2 p_i = p; prec \sum C_i$	Coffman & Graham (1972)
$Pm p_i = p; tree \sum C_i$	Baptiste et al. (2004)
$Qm p_i = p; r_i \sum C_i$	Dessouky et al. (1990)
$R \sum C_i$	Born (1973), Bruno et al. (1974)
$P p_i = p; r_i \sum w_i C_i$	Brucker & Kravchenko (2008)
$P p_i = 1; r_i \sum w_i U_i$	Networkflowproblem
$Pm p_i = p; r_i \sum w_i U_i$	Baptiste et al. (2004)
$Q p_i = p \sum w_i U_i$	Assignment-problem
$P p_i = p; r_i \sum T_i$	Brucker & Kravchenko (2005)
$P p_i = 1; r_i \sum w_i T_i$	Networkflowproblem
$Q p_i = p \sum w_i T_i$	Assignment-problem

• maximal pseudopolynomially solvable:

$Qm r_i C_{max}$	Lawler et al. (1989)
$Qm \sum w_i C_i$	Lawler et al. (1989)
$Qm \sum w_i U_i$	Lawler et al. (1989)

• minimal NP-hard:

$P2 C_{max}$	Lenstra et al. (1977)
$* P C_{max}$	Garey & Johnson (1978)
$* P p_i = 1; intree; r_i C_{max}$	Brucker et al. (1977)
$* P p_i = 1; prec C_{max}$	Ullman (1975)
$* P2 chains C_{max}$	Du et al. (1991)
$* Q p_i = p; chains C_{max}$	Kubiak (1988)
$* P p_i = 1; outtree L_{max}$	Brucker et al. (1977)
$* P p_i = 1; intree; r_i \sum C_i$	Lenstra (-)
$* P p_i = 1; prec \sum C_i$	Lenstra & Rinnooy Kan (1978)

Graham's List Scheduling for Minimizing Makespan ('66)


Algorithm 1 LS

- 1: sort jobs arbitrarily, J_1, \dots, J_n
 - 2: **for** $j = 1, \dots, n$ **do**
 - 3: assign J_j to currently least loaded machine
 - 4: **end for**
-

- Probably the first approximation algorithm ever
- Actually an *online* algorithm
 - since order is arbitrary
- How bad can LS be?
 - Assume job order is (considering processing times)

$$\underbrace{1, 1, \dots, 1}_{m(m-1) \text{ jobs}}, m$$

- $\frac{\text{LS}}{\text{OPT}} = \frac{2m-1}{m} = 2 - \frac{1}{m}$



Competitive Analysis of List Scheduling

Theorem

Algorithm LS is $(2 - \frac{1}{m})$ -approximation (competitive)

Proof.

- Observation: $\text{OPT} \geq \max \left\{ \frac{\sum_j p_j}{m}, \max_j p_j \right\}$
- Notation:
 - L : maximum load in LS
 - i^* : most loaded machine in LS
 - j^* : last job assigned to i^*
- All machines have load $\geq L - p_{j^*}$
 $\Rightarrow (\sum_j p_j) - p_{j^*} \geq m(L - p_{j^*})$

$$\begin{aligned} L &\leq \frac{(\sum_j p_j) - p_{j^*}}{m} + p_{j^*} \\ &= \frac{\sum_j p_j}{m} + \left(1 - \frac{1}{m}\right) p_{j^*} \\ &\leq \left(2 - \frac{1}{m}\right) \text{OPT} \end{aligned}$$



Improving upon List Scheduling

- Observations:
 - if $n \leq m$, LS is optimal
 - ensuring load-balancing online might be harmful
- How would improved online algorithms look like?
 - purposely cause imbalance
 - prepare for the worse
 - leave “some” room for a big job
 - not too much

	deterministic			randomized	
m	lower bound	upper bound	LS	lower bound	upper bound
2	1.5000	1.5000	1.5000	1.3333	1.3334
3	1.6666	1.6667	1.6667	1.4210	1.5567
4	1.7310	1.7333	1.7500	1.4628	1.6589
5	1.7462	1.7708	1.8000	1.4873	1.7338
6	1.7730	1.8000	1.8333	1.5035	1.7829
7	1.7910	1.8229	1.8571	1.5149	1.8169
∞	1.8520	1.9230	2.0000	1.5819	—

Improving upon List Scheduling

- What if we're allowed to sort jobs?
 - no longer online...
 - can we get a better approximation guarantee?
- ORDEREDLS: jobs are sorted: $p_1 \geq p_2 \geq \dots \geq p_n$

Lemma

If $n > m$ then $\text{OPT} \geq p_m + p_{m+1}$

Proof.

- Pigeonhole principle (for any solution):
 - 2 jobs out of J_1, \dots, J_{m+1} are assigned to same machine
 - jobs are ordered by non-increasing processing time
 - the minimum weight of any such 2 jobs is $p_m + p_{m+1}$



Analysis of Improved Offline List Scheduling

Theorem

Algorithm ORDEREDLS is a $(\frac{3}{2} - \frac{1}{2m})$ -approximation

Proof (similar to previous).

- Observation: $\text{OPT} \geq \max \left\{ \frac{\sum_j p_j}{m}, \max_j p_j \right\}$
- Same notation:
 - L (max-load), i^* (max-load machine), j^* (last job in i^*)
- If $j^* \leq m$: $\text{ORDEREDLS} = \text{OPT}$
- If $j^* > m$: $p_{j^*} \leq \frac{p_m + p_{m+1}}{2} \leq \frac{\text{OPT}}{2}$

$$\begin{aligned}
 L &\leq \frac{(\sum_j p_j) - p_{j^*}}{m} + p_{j^*} \\
 &= \frac{\sum_j p_j}{m} + \left(1 - \frac{1}{m}\right) p_{j^*} \\
 &\leq \text{OPT} + \left(1 - \frac{1}{m}\right) \frac{\text{OPT}}{2} \\
 &\leq \left(\frac{3}{2} - \frac{1}{2m}\right) \text{OPT}
 \end{aligned}$$



Interval Scheduling

- Input
 - job J_j requires processing during interval $I_j = [s_j, t_j)$
 - uses *all* the interval
 - each job J_j has weight $w(J_j)$
 - m identical machines
- Goal
 - find a subset of jobs $S = S_1 \dot{\cup} S_2 \dot{\cup} \dots \dot{\cup} S_m$, such that
 - for any k , and any $J_j, J_{j'} \in S_k$, $I_j \cap I_{j'} = \emptyset$
 - criteria
 - $\max \sum_{J_j \in S} w(J_j)$
- Usually referred to as *interval scheduling*
 - refer to J_j as its interval
- Our focus:
 - uniform weights, i.e., $w(J_j) = 1$ for all j
- What algorithm would you use?

Offline Interval Scheduling ($m = 1$)

Algorithm 2 GIS(machine, jobs \mathcal{J})

```
1: sort intervals in  $\mathcal{J}$  ▷ How?
2: for  $j = 1, \dots, n$  do
3:   if  $J_j$  doesn't overlap any currently scheduled interval then
4:     add  $J_j$ 
5:   end if
6: end for
```

- Greedy interval scheduling (GIS...)
- Main question: how should we sort intervals?

Offline Interval Scheduling ($m = 1$)

Theorem

GIS (sorting by t_j) produces an optimal schedule

Proof (by contradiction).

- Let J_{j_1}, \dots, J_{j_k} be the schedule produced by GIS
- Let OPT be an optimal schedule
 - OPT schedules $J_{j_1}, \dots, J_{j_r}, J_{j_{r+1}^*}, \dots$
 - $r \leq k$ is maximal
 - i.e., OPT “agrees” with GIS on a maximum prefix of GIS
- if $r = k$, GIS would have scheduled additional jobs
 - at least $J_{j_{r+1}^*}$ is available and non-overlapping
- Assume $r < k$
 - by definition of GIS, $t_{j_{r+1}} \leq t_{j_{r+1}^*}$
 - we can switch them without harming OPT
- contradiction to GIS definition / maximality of r



Offline Interval Scheduling ($m \geq 1$)

Algorithm 3 m -GIS(m machines, jobs \mathcal{J})

- 1: **for** $i = 1, \dots, m$ **do**
 - 2: add GIS(machine i , remaining jobs) \triangleright GIS $_i$
 - 3: **end for**
-

Theorem

m -GIS is a 2-approximation

Proof (by contradiction).

- OPT_i : set of intervals in $\text{OPT} \setminus m\text{-GIS}$ scheduled on M_i
- Assume $|m\text{-GIS}| < |\text{OPT}|/2$
 $\Rightarrow |\cup_i \text{OPT}_i| > |\text{OPT}|/2 > |m\text{-GIS}| = |\cup_i \text{GIS}_i|$
- Pigeonhole principle
 - exists i s.t. $|\text{OPT}_i| > |\text{GIS}_i|$
 - OPT_i is available to GIS_i : contradicts optimality of GIS \square

Online Interval Scheduling

- Job J_j arrives at s_j
 - reveals t_j
- preemptive
 - otherwise, same as ordering by s_j (unbounded competitiveness)

Algorithm 4 ONLINE- m -GIS: at time t

```
1: for any job  $J_j$  arriving at  $t$  do
2:   if  $J_j$  can be scheduled on some machine  $M_i$  then
3:     schedule  $J_j$  on  $M_i$ 
4:   else
5:      $j' \leftarrow \arg \max_k \{t_k \mid J_k \text{ is scheduled and intersects } J_j\}$ 
6:     if  $t_j \geq t_{j'}$  then drop  $J_j$ 
7:     else drop  $J_{j'}$  and schedule  $J_j$  instead
8:     end if
9:   end if
10: end for
```

Online Interval Scheduling

- ONLINE- m -GIS emulates m -GIS

Corollary

ONLINE- m -GIS is 2-competitive

- Is this tight?
- If m -GIS is an α -approximation
 - ONLINE- m -GIS is α -competitive

Theorem

m -GIS produces an optimal schedule

Corollary

ONLINE- m -GIS produces an optimal schedule

- Online!!

From Networking to Scheduling

- Links / paths \longrightarrow machines
 - Traffic \longrightarrow jobs
 - Applications:
 - routing
 - optical networks
 - Wavelength Division Multiplexing (WDM)
 - packet scheduling
 - bounded-delay model
 - forwarding across networks (sensors, bounded-buffers)
 - many more...
-
- Some references
 - Cormen, Leiserson, Rivest. Introduction to Algorithms, MIT Press, Cambridge, MA, 1990.
 - Sgall, Online Scheduling - A Survey, in Fiat et al (eds.), Online Algorithms: The State of the Art, 1998.
 - Graham et al., Optimization and Approximation in Deterministic Sequencing and Scheduling: a Survey, 1979.

The Creative Part of the Course

- Models we've seen:
 - AAP Routing
 - BM with commitments (FIFO)
 - BM with packet weights (FIFO)
 - BM/Scheduling with packet weights and deadlines
- Issues to take into account:
 - interesting/useful model extension
 - better algorithms
 - proofs (analytic)
 - simulations (heuristic)
 - KISS: "Keep It Simple, Stup..."
- Analytic approach:
 - toy examples are usually very instructive
 - limited models are a good place to start
- Simulation approach:
 - try to highlight the main ideas/benefits
 - a great motivator for designing better solutions