

# Buffer Management and Scheduling with Packet Dependencies


Gabriel Scalosub, Ben Gurion University

Based on joint works with:

Alex Kesselman, Boaz Patt-Shamir, Peter Marbach, and Jörg Liebeherr

# Motivation: Video Streaming

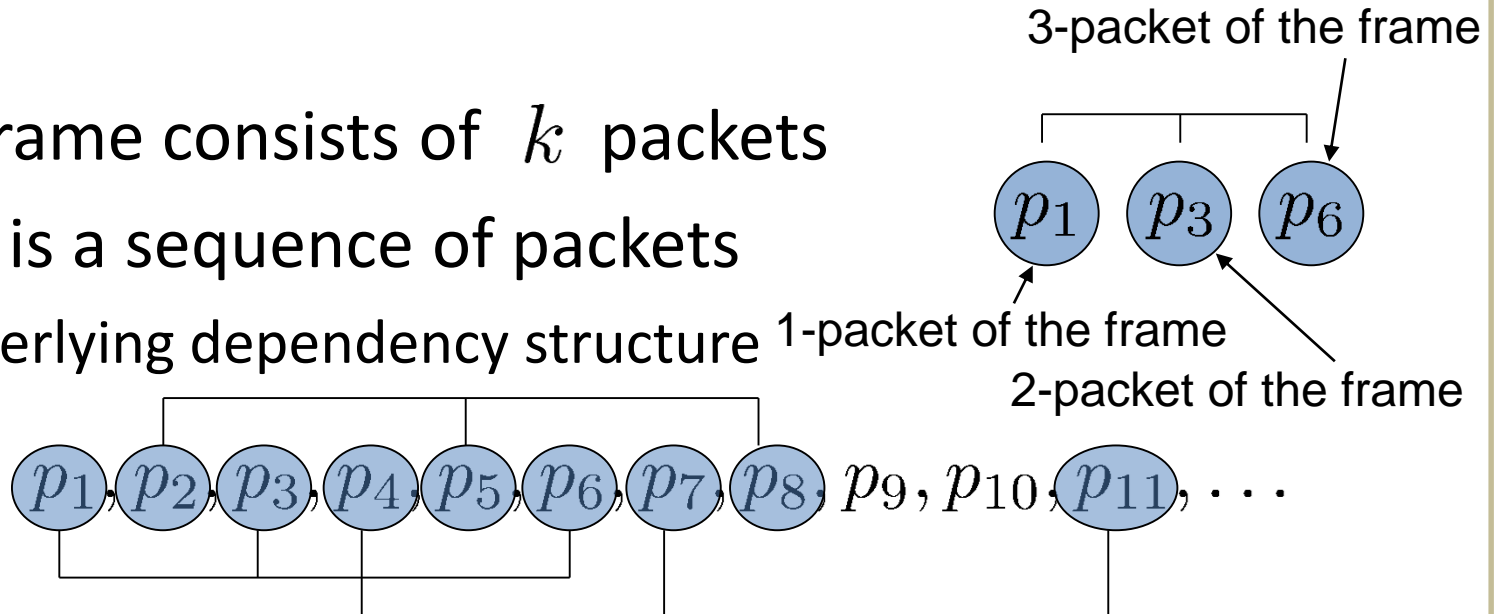
- Smart encoding:
  - Suffices to recover many



- Every video frame is fragmented into packets
- Restoration depends on recovering all packets
- If packets are lost:
  - Affects other packets as well (become redundant)
  - Streaming: retransmission is not an option
    - Popularity: Live TV over IP constantly increasing

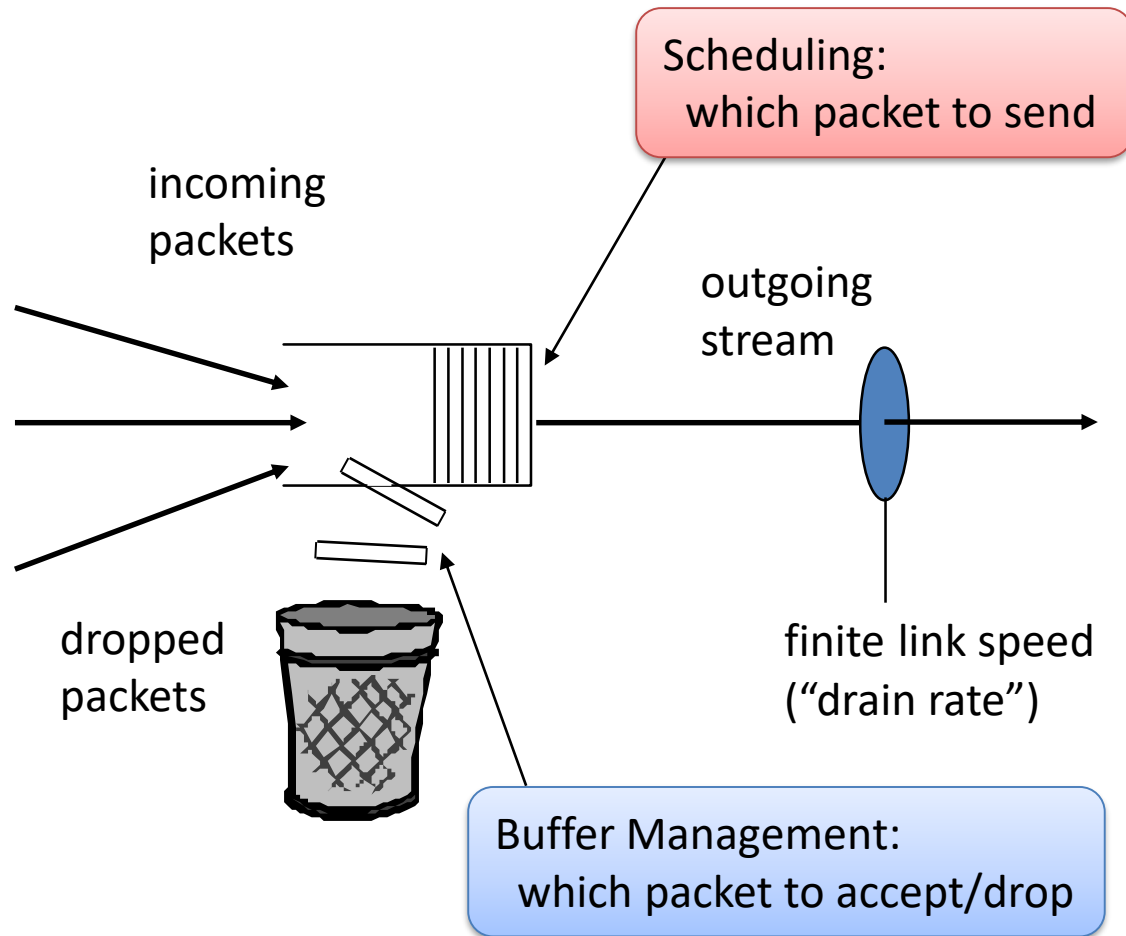
# Traffic, Dependencies, and Goodput

- Each frame consists of  $k$  packets
- Traffic is a sequence of packets
  - Underlying dependency structure



- Goodput: number of whole frames delivered
- Throughput vs. Goodput
  - Throughput (packet-level)  $\neq$  Goodput (frame-level)
  - Also verified by experimental studies (e.g., MPEG)

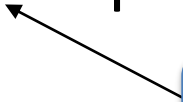
# Buffering Schematics



# Buffer Model

- Single FIFO queue of size  $B$

“fix” scheduling,  
design buffer management



- Discrete time:

- Delivery substep

- One packet delivered from head of queue

- Arrival substep

- Packets arrive
    - Some packets may be dropped
    - Packets accommodated in the buffer

# Methodology

- Adversarial traffic
  - An adversary generates the traffic dynamically
  - “Provides” an optimal solution for this traffic
  - Analysis globally applicable
    - Independent of the process generating the traffic
- Assumption:
  - An online algorithm has no knowledge of future traffic
  - OPT is the optimal clairvoyant (offline) algorithm
- Algorithm  $A$  is  $c$ -competitive if  $\forall$  finite traffic  $\sigma$ 
$$A(\sigma) \geq \frac{1}{c} \cdot \text{OPT}(\sigma)$$

# The Problem

- Given:
  - Finite size buffer
  - FIFO scheduler
  - Incoming traffic with packet dependencies
- Goal:
  - Maximize the goodput
- Main difficulty:
  - What to do upon overflow?? Which packet(s) to drop?

# Our Contributions

- Design guidelines for algorithms
  - *No-regret*:  
Once a packet is admitted to the buffer (not necessarily sent), make every attempt possible to deliver the frame.
  - *Ensure progress*:  
Deliver a complete frame as soon as possible
- Competitive algorithms (following these guidelines)
  - Analytic guarantees
  - Good performance in simulations
- Lower bounds




# Related Work

- MPEG inter-frame dependencies
  - Consider frames as “basic” entities (although routers decide on packets)  
[Zhang et al., 2001], [Awad et al., 2002]
- Inter-packet dependencies
  - Implementation [Ramanathan et al., 1995]
  - ATM guaranteed frame rate [Bonaventure & Nelissen, 2001]
- Buffer management for QoS
  - Multi-valued packets, competitive analysis  
[Lapid et al., 2000], [Kesselman et al., 2004], [Englert & Westerman, 2006]
- Proactive coding and FEC [Albanese et al., 1996]


# Preliminaries

- Offline
  - Closely related to  $k$ -DM ( $\Omega(k/\log k)$ -hard to approx.)
  - Simple greedy algorithm is a  $(k+1)$ -approximation
- Online (arbitrary traffic)
  - Not much you can do



*ALG*

*OPT*



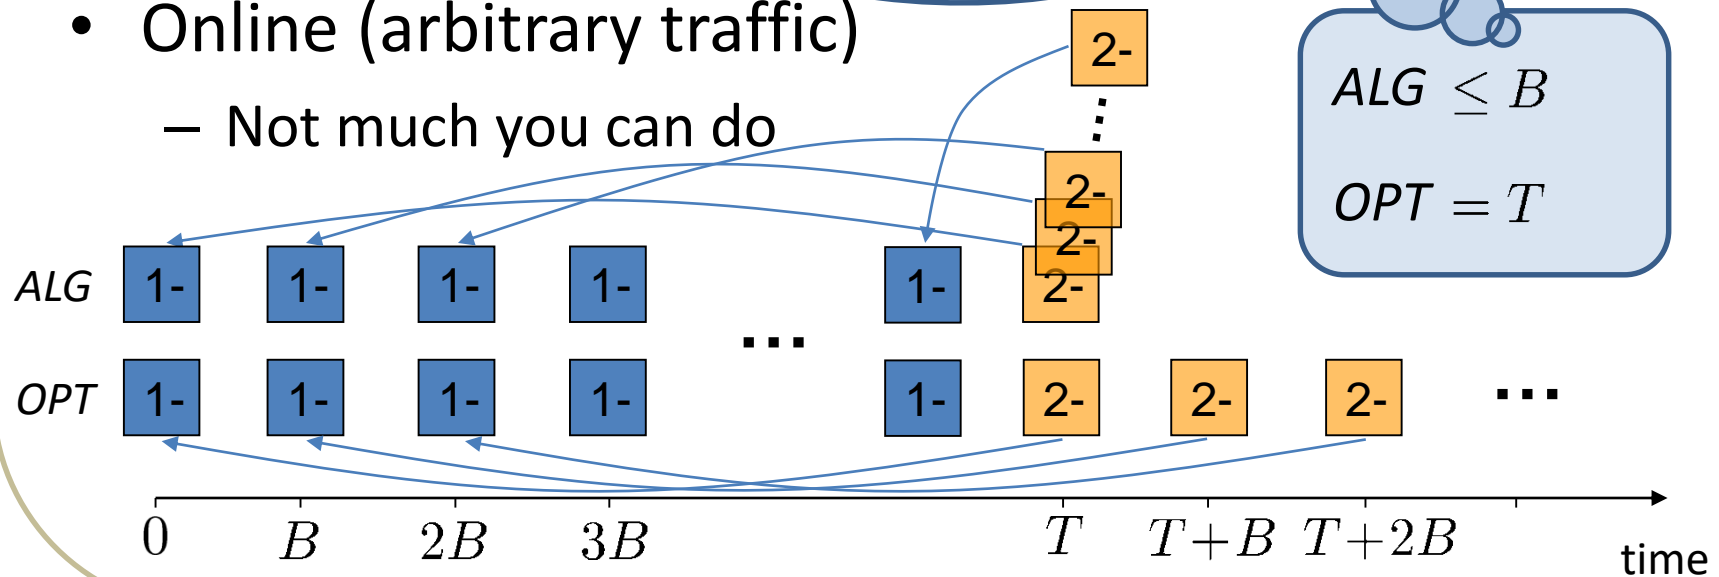
# Preliminaries

Some thoughts:

- (1) Is this reasonable traffic?
- (2) Maybe the adversary is too strong??

- Online (arbitrary traffic)

- Not much you can do



# Order Respecting Traffic

- “Problem”:
  - Selective unbounded delay/burstiness
- Model requirement (“solution”):
  - Both ALG and OPT have to deal with same delay/burstiness
- Order-respecting traffic:
  - Frame order induced by  $j$ -packets is the same for every  $j$ 
    - OK: 2.1 (frame 2, part 1), 3.1, 2.2, 3.2
    - Not OK: 3.1, 2.1, 2.2, 3.2
- Lower bound for order respecting traffic:  $\Omega(k)$

# A Greedy Algorithm (*GA*)

- Consider  $k=2$

At every time  $t$  where overflow occurs, keep packets according to the following priority:

1. 2-packets whose 1-packets were delivered
2. Complete frames (both 1-packets and 2-packets)
3. Remaining 1-packets

- Theorem: *GA* is constant-competitive for  $k=2$
- What if  $k>3$ ?

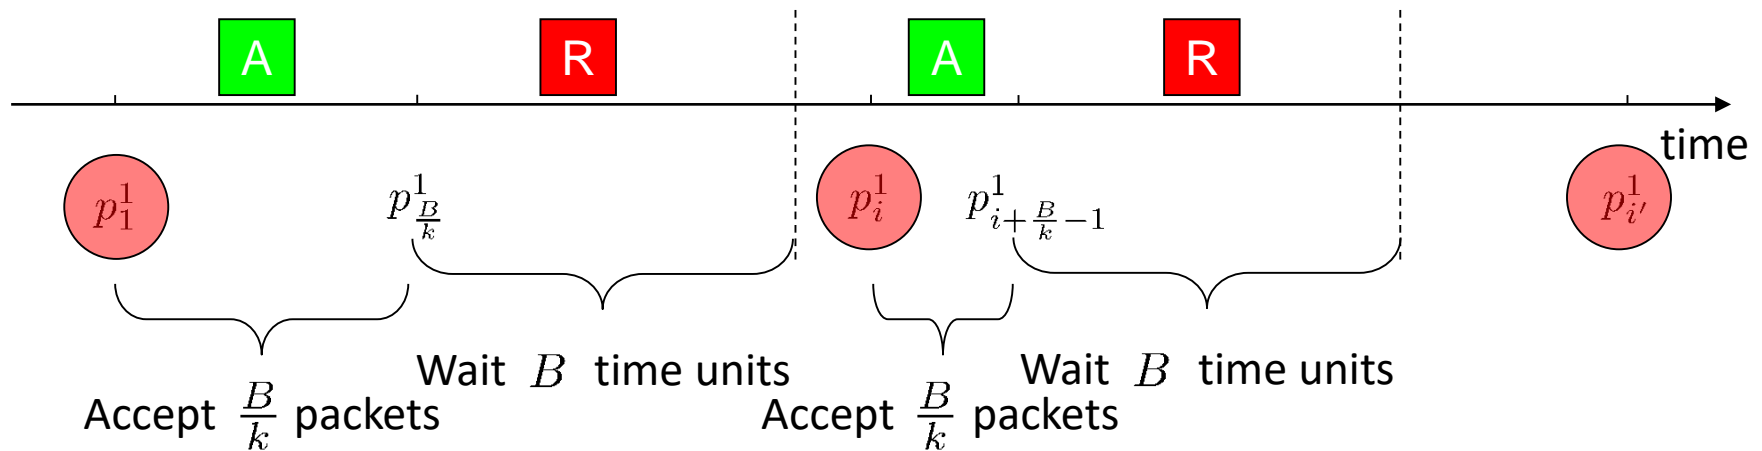
# Static-Partitioning Algorithm (*SPA*)

- Intuition
  - Think ahead: focus on packet admission
  - Virtually partition the buffer into  $k$  levels of size  $\frac{B}{k}$ 
    - Buffer is still FIFO!!
  - Level  $j$  only holds  $j$ -packets
  - Level  $j$  accepts  $j$ -packets that are “evenly” spaced in time
    - Alternating accept/reject periods
  - Levels synchronize on frame index
    - Ensures delivered packets correspond to the same frame
- Extra perk: non-preemptive

# Example: SPA for $k=2$

$p_i^j$  : Frame  $f_i$ 's  
 $j$ -packet

- Consider level 1, i.e., 1-packets




$p_i^1$  is the first 1-packet arriving after reject period

# Example: SPA for $k=2$

$p_i^j$  : Frame  $f_i$ 's  
 $j$ -packet

- Consider level 1, i.e., 1-packets



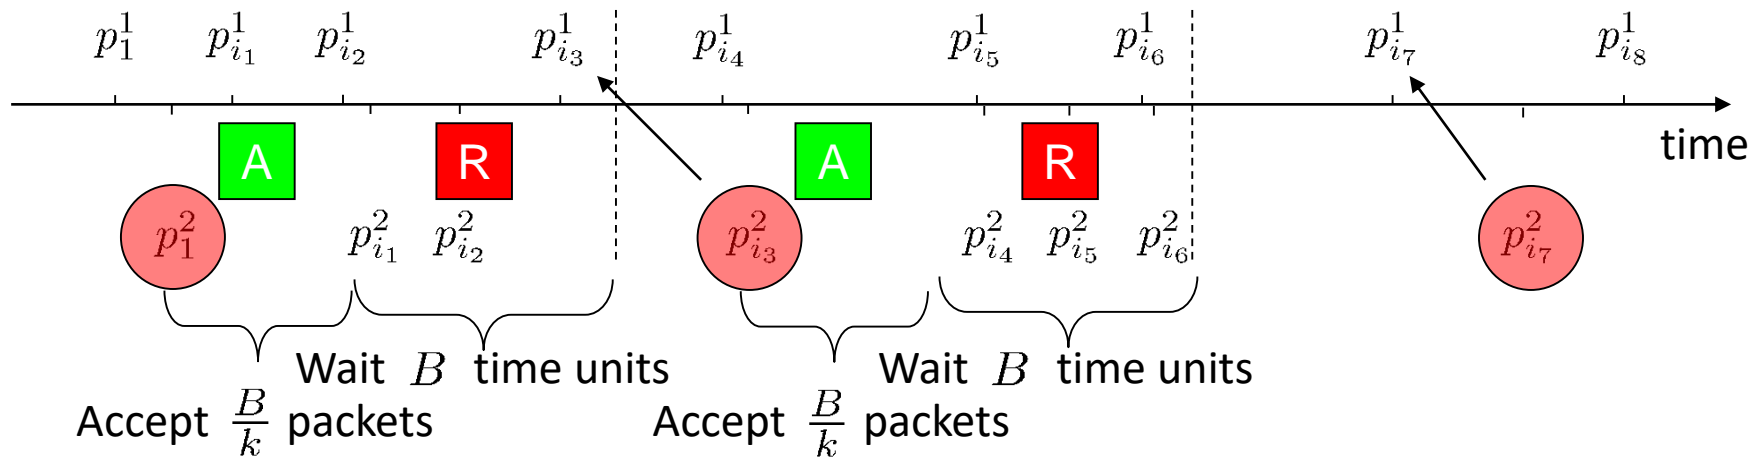
- 1-sync frame indices:  $1, i, i', \dots$
- Accepts  $\frac{B}{k}$  first 1-packets after every 1-sync
  - Specifically, has sufficient buffer space



# Example: SPA for $k=2$

$p_i^j$  : Frame  $f_i$ 's  
 $j$ -packet

- Consider level 2, i.e., 2-packets




$p_{i_3}^2$  is the first 2-packet of a 1-sync arriving after reject period

# Example: SPA for $k=2$

$p_i^j$  : Frame  $f_i$ 's  
 $j$ -packet

- Consider level 2, i.e., 2-packets




- 2-sync indices  $\subseteq$  1-sync indices
- Accepts  $\frac{B}{k}$  first 2-packets after every 2-sync
  - Specifically, has sufficient buffer space

# Analysis of *SPA*

- Analysis ingredients:
  - For any two consecutive  $k$ -syncs  $i, i'$   
*SPA* delivers  $\frac{B}{k}$  out of frames  $\{f_i, \dots, f_{i'-1}\}$
  - For any two consecutive  $k$ -syncs  $i, i'$   
*OPT* delivers at most  $2kB + B$  out of frames  $\{f_i, \dots, f_{i'-1}\}$ 
    - Intuition: at most  $\frac{B}{k} + 2B$  out of every level
- Theorem: *SPA* is  $(2k^2 + k)$ -competitive
- Design criteria:
  - *No regret*: algorithm never preempts
  - *Ensure progress*: the first  $\frac{B}{k}$  frames are always delivered

# Coming Closer to Reality

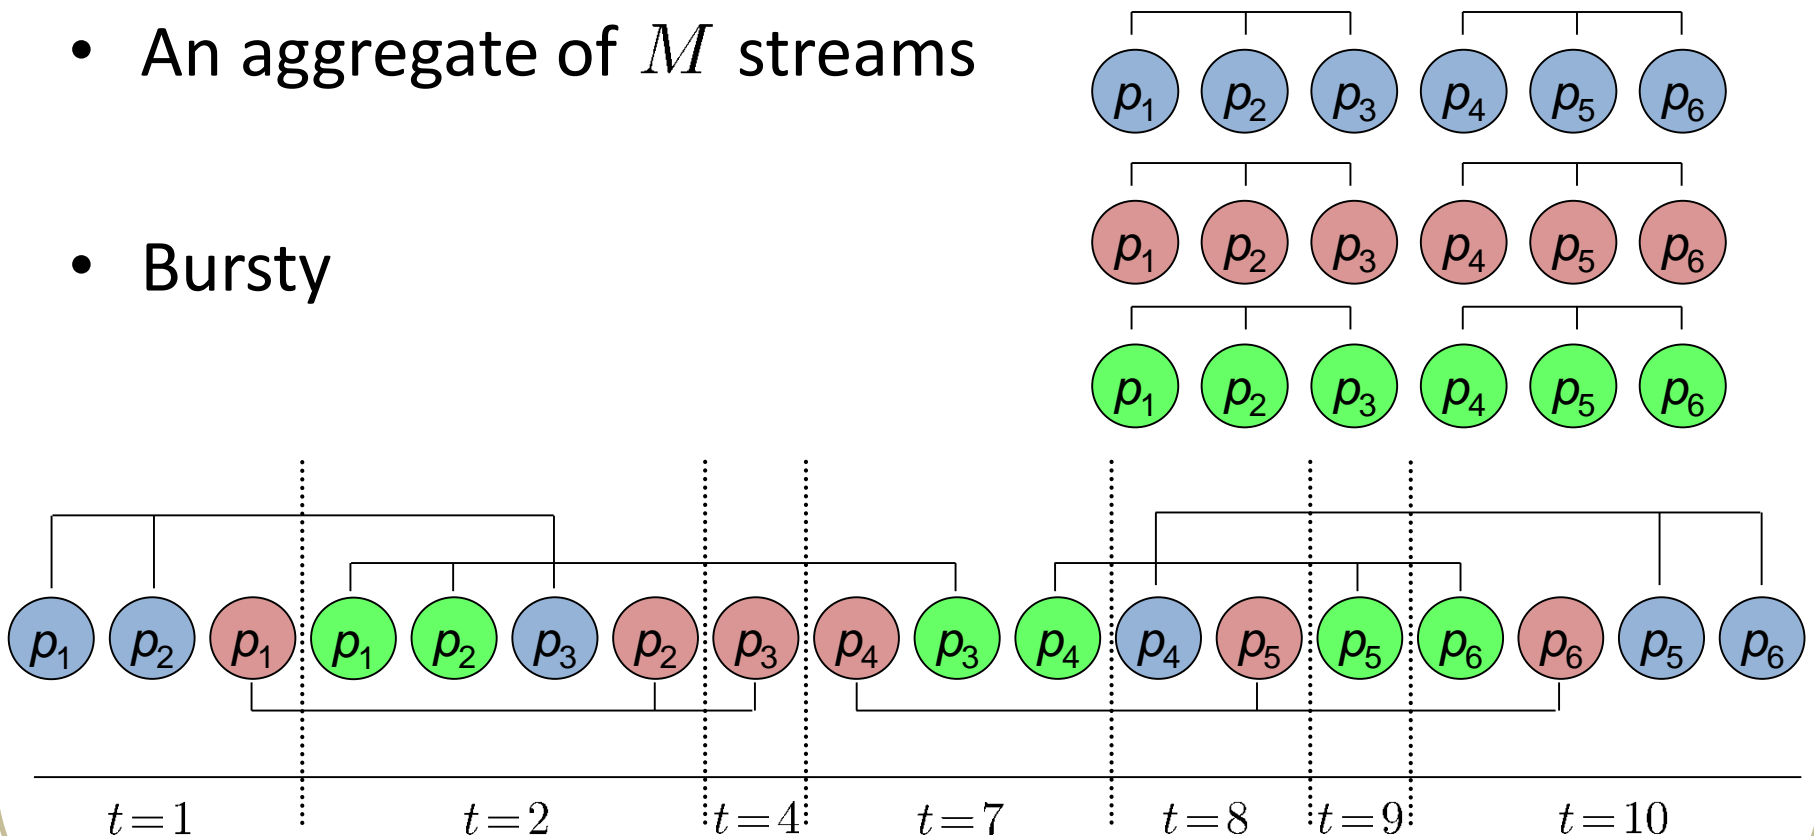
- Real-life traffic is hardly order-respecting
  - Network traffic consists of multiple *flows/streams*
  - No reason to assume any order is maintained for frames of different streams
- Modeling a single stream:



- Packet stream is partitioned into blocks of  $k$  packets
- Each block is a single frame
- Packets of a stream arrive in order

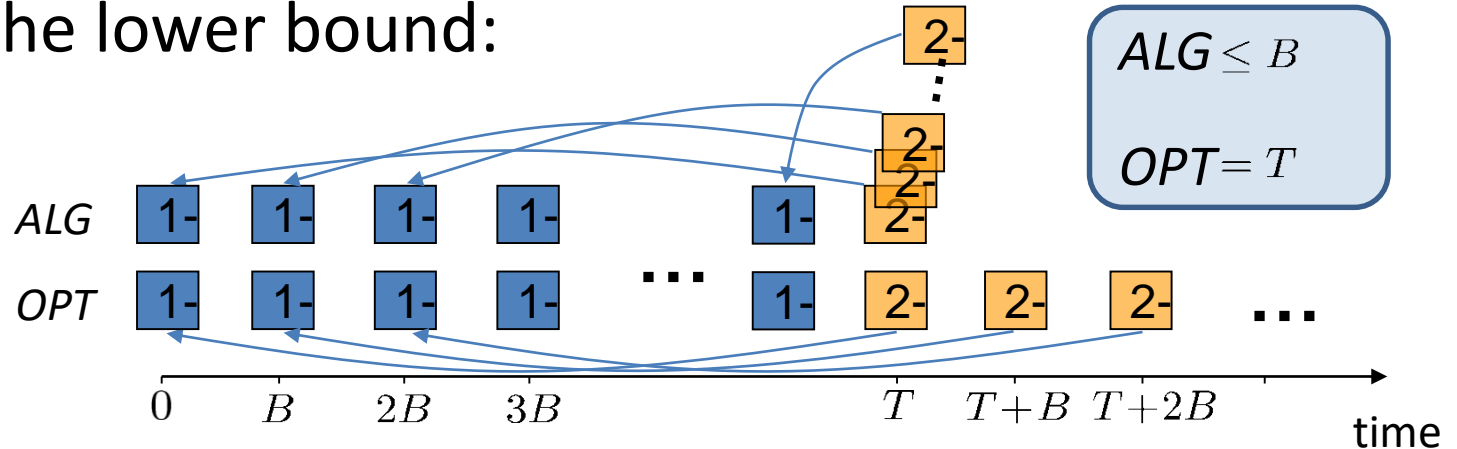
# Multiple Streams

- An aggregate of  $M$  streams
- Bursty



# Lower Bounds

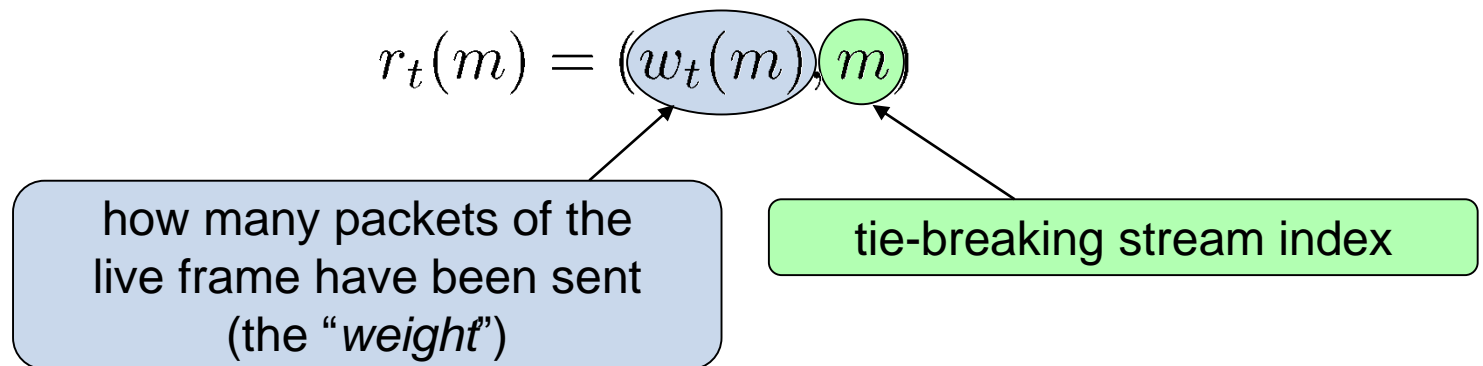
- Recall the lower bound:



- $M$  streams implies  $T = O(M)$ 
  - Immediately gives a lower bound of  $\Omega\left(\frac{M}{B}\right)$
- Can be generalized to show a lower bound of  $\Omega\left(\frac{kM}{B}\right)$

# The Return of the Greedy Approach


- At any time  $t$ :
  - Every stream has at most one “live” frame
  - We define a ranking on all streams by (lexicographically):



- Algorithm *WeightPriority* (*WP*)
  - Upon overflow, drop lowest ranking frames first

# Analysis of $WP$ (in a nutshell)


- Map frames delivered by  $OPT$  to ones delivered by  $WP$ 
  - Partition time into intervals  $I_1, I_2, I_3, \dots$
  - Map every overflow in  $I_j$  to a packet delivered by  $WP$  in  $I_j$ 
    - Implicitly implies a maximum weight for every interval






# Analysis of $WP$ (in a nutshell)

- Map frames delivered by  $OPT$  to ones delivered by  $WP$ 
  - Partition time into intervals  $I_1, I_2, I_3, \dots$
  - Map every overflow in  $I_j$  to a packet delivered by  $WP$  in  $I_j$ 
    - Implicitly implies a maximum weight for every interval
  - Map every  $I_j$  to a strictly-higher weight interval



# Analysis of $WP$ (in a nutshell)

- Show there exists an  $\alpha$ :
  - $OPT$  handles at most  $\alpha$  frames in each interval  $I_j$
  - At most  $\alpha$  intervals are mapped to any interval  $I_j$
- Results in a  $k$ -height  $\alpha$ -ary tree
- Choosing  $\alpha \sim kMB$  turns out to be a good choice
- Theorem:  $WP$  is  $O((2kMB)^{k+1})$ -competitive.



# Analysis of $WP$ (in a nutshell)

- Show there exists an  $\alpha$ :
  - $OPT$  handles at most  $\alpha$  frames in each interval  $I_j$
  - At most  $\alpha$  intervals are mapped to any interval  $I_j$
- Results in a  $k$ -height  $\alpha$ -ary tree
- Choosing  $\alpha \sim kMB$  turns out to be a good choice
- Theorem:  $WP$  is  $O((2kMB)^{k+1})$ -competitive.
- Design criteria: due to ranking

$$r_t(m) = (w_t(m), m)$$

Ensure progress


No regret

# Additional Algorithms

- *FrameOblivious*
  - drop overflowing packets (ignore dependencies)
  - no state information req., no buffer scan
- *SemiFrameOblivious*
  - Same as *FrameOblivious*, but
  - Scans buffer to remove packets of dropped frames
- *StreamOblivious*
  - Same as *WP*, but rank is solely based on weight (arbitrary tie-break)
- *StreamPriority*
  - Same as *WP*, but rank is solely based on stream index

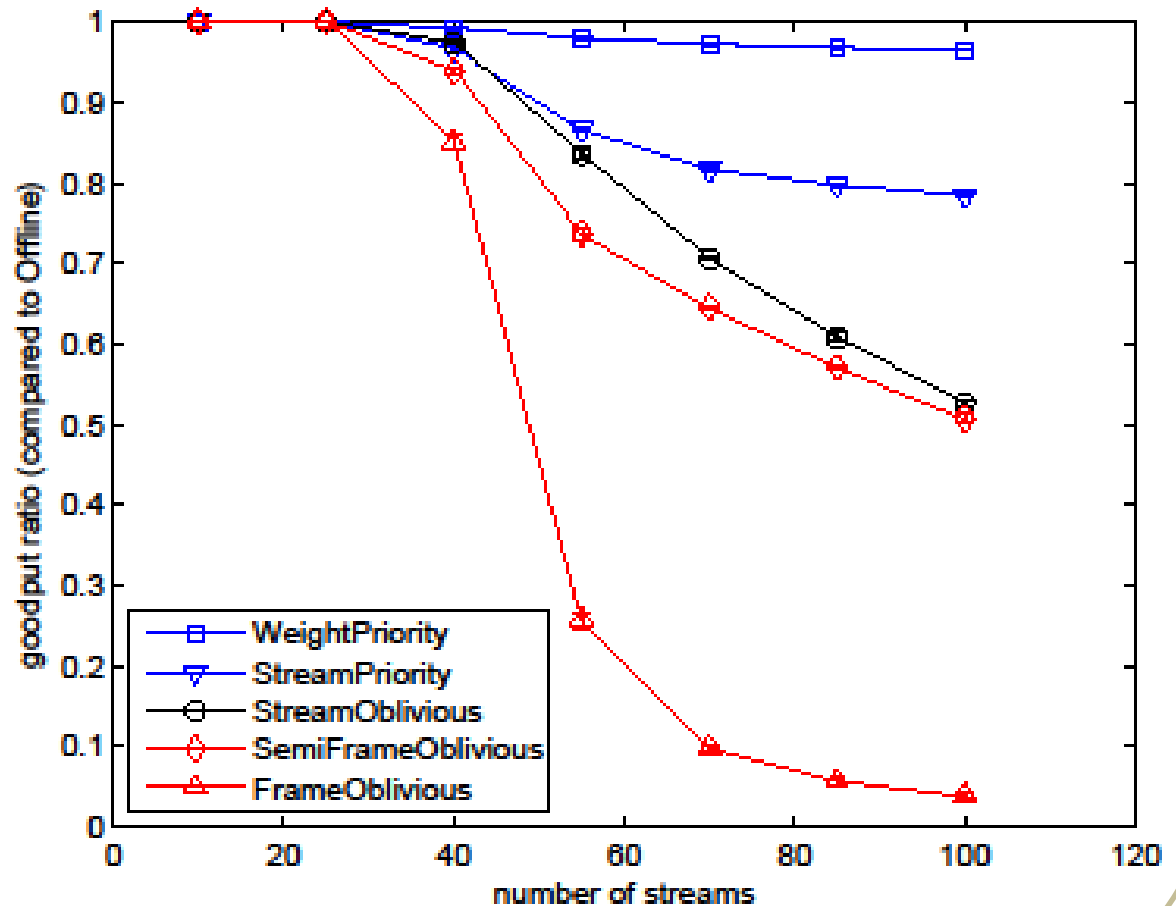
# Effect of increasing $k$

- Each stream:
  - MMPP source
  - Av. rate: 0.025
- 50 streams
  - Av. agg. rate: 1.25
- 200 frames/stream
- $B=12$



# Effect of increasing $M$

- Each stream:
  - MMPP source
  - Av. rate: 0.025
- 200 frames/stream
- $B=12, k=6$



# Summary

- New model for traffic with packet dependencies
- Competitive algorithms (and lower bounds):

Traffic Model	Lower Bound	Upper Bound	Conclusion
Arbitrary	Unbounded		Too General
Order-resp.	$\Omega(k)$	$O(k^2)$	Too restrictive
$M$ Streams	$\Omega\left(\frac{kM}{B}\right)$	$O((2kMB)^{k+1})$	Manageable, Realistic

Later improved to  $O(k)$

Later improved to  $O(kMB)$  (deterministic) and  $O(M)$  (randomized)

# Additional Models and Questions

- Competitive ratio gaps (new algorithms?)
- Using forward-error-correction (FEC)
  - Suffices to deliver  $d$ -of- $k$
- Recouple scheduling
  - “Price of FIFOness”
- Other extensions and QoS considerations
  - Fairness, increasing/diminishing returns, variable frame size / value, **delay**
- Many more
  - Inter-frame dependencies, randomness, ...

Most results are for the case of no buffers:

a single packet admitted and delivered in every time slot

Focus on *scheduling*, some results on uniform delays



# References

- A. Kesselman, B. Patt-Shamir and G. Scalosub, “Competitive Buffer Management with Packet Dependencies”, Theoretical Computer Science, 489--490, pp. 75-87, 2013.
- G. Scalosub, J. Liebeherr and P. Marbach, “Buffer Management for Aggregated Streaming Data with Packet Dependencies”, IEEE Transactions on Parallel and Distributed Systems 24(3), pp. 439-449, 2013.
- Yishay Mansour, Boaz Patt-Shamir, and Dror Rawitz, “Overflow management with multipart packets”, Computer Networks, 56(15), pp. 3456–3467, 2012.
- Y. Emek, M. M. Halldórsson, Y. Mansour, B. Patt-Shamir, J. Radhakrishnan, and D. Rawitz, “Online set packing and competitive scheduling of multi-part tasks”, SIAM Journal on Computing, 41(4), pp. 728–746, 2012.
- J. Kawahara, K. M. Kobayashi, and S. Miyazaki, “Better bounds for online k-frame throughput maximization in network switches”, TCS, 657(B), pp. 173-190, 2017.
- G. Scalosub, “Towards Optimal Buffer Management for Streams with Packet Dependencies”, Computer Networks, 129, pp. 207-214, 2017.
- M. Markovitch and G. Scalosub, “Bounded Delay Scheduling with Packet Dependencies”. Computer Communications, 182, pp. 98-109, 2022.

# Network QoS

## 371-2-0213

### Lecture 13

Gabriel Scalosub

# Outline

- 1 Postscript
  - Context of Our Course
  - QoS Research Angles

# Course Agenda


- Design and analysis of protocols for providing *Quality-of-Service* (QoS):
  - currently used protocols, and
  - new designs and models
- Main problems in networks arise from drops/delays, e.g.:
  - applications: bumpy audio/video
  - network: increased congestion due to TCP retransmissions
- Network resources (BW, buffers, ...) shared by all the users
- Most current networks: best-effort
  - all packets were created equal
  - FIFO scheduling
  - no admission control
- QoS support must be active about resource allocation
  - how should resources be allocated?

# Current Concepts of QoS

- Integrated Services (IntServ)
  - per-flow resource reservation
    - routing protocol finds E2E path depending on flow demand
    - protocol: RSVP
  - handicaps:
    - accounting (e.g., inter-AS), scalability (per-flow state)
- Differentiated Services (DiffServ)
  - better than best-effort
  - dividing traffic into small number of forwarding classes
    - employ policing, marking per class
    - traffic prioritization: scheduling/buffer management at routers
  - handicaps:
    - no reservation, hard to provide deterministic QoS guarantees
- Traffic engineering
  - optimize usage of network resources
  - sophisticated routing
    - needs “global” information: topology, traffic (mostly intra-AS)
    - usually MPLS-based

# Routing and Buffer Management


- AAP routing
  - competitive model for path reservation
  - models many of IntServ's aspects
- Committed and excess buffer mgmt.
  - competitive model for satisfying SLA
  - models aspects of managing a single AF PHB class



Feasibility and lag of algorithm

# Routing and Buffer Management


- AAP routing
  - competitive model for path reservation
  - models many of IntServ's aspects
- Committed and excess buffer mgmt.
  - competitive model for satisfying SLA
  - models aspects of managing a single AF PHB class
- Weight-based FIFO buffer mgmt.
  - competitive model for managing multiple AF/EF PHBs



Analysis of Greedy is tight

# Routing and Buffer Management

- AAP routing
  - competitive model for path reservation
  - models many of IntServ's aspects
- Committed and excess buffer mgmt.
  - competitive model for satisfying SLA
  - models aspects of managing a single AF PHB class
- Weight-based FIFO buffer mgmt.
  - competitive model for managing multiple AF/EF PHBs
- Bounded-delay scheduling and buffer mgmt.
  - TTL-based competitive model for managing delay-sensitive traffic




Asymptotic LB of  $\phi \sim 1.618$




# Routing and Buffer Management

- AAP routing
  - competitive model for path reservation
  - models many of IntServ's aspects
- Committed and excess buffer mgmt.
  - competitive model for satisfying SLA
  - models aspects of managing a single AF PHB class
- Weight-based FIFO buffer mgmt.
  - competitive model for managing multiple AF/EF PHBs
- Bounded-delay scheduling and buffer mgmt.
  - TTL-based competitive model for managing delay-sensitive traffic
- Packet dependencies
  - packet level decisions vs. data-frame level effect



# Network Calculus and Scheduling

- Network Calculus
  - Arrival and service curves
  - Envelopes and regulators
  - Prove IntServ delay bounds



Arrival and service curves

# Network Calculus and Scheduling


- Network Calculus
  - Arrival and service curves
  - Envelopes and regulators
  - Prove IntServ delay bounds
- Classic machine scheduling
  - Makespan, interval scheduling



bad ways to sort intervals


# Network Calculus and Scheduling

- Network Calculus
  - Arrival and service curves
  - Envelopes and regulators
  - Prove IntServ delay bounds
- Classic machine scheduling
  - Makespan, interval scheduling
- Cornerstones of packet scheduling
  - FIFO, PQ, RR, WFQ (and beyond)



# Network Calculus and Scheduling

- Network Calculus
  - Arrival and service curves
  - Envelopes and regulators
  - Prove IntServ delay bounds
- Classic machine scheduling
  - Makespan, interval scheduling
- Cornerstones of packet scheduling
  - FIFO, PQ, RR, WFQ (and beyond)
- Adversarial Queueing Theory (AQT)
  - scheduling across a network
  - stability of protocols, traffic-patterns, and topologies



AQT arrival model

# Many Interesting Research Directions

- Buffer management
    - e.g., packet dependencies, closing UB/LB gaps, other models
  - Scheduling
    - interplay between scheduling and buffer management
    - it shouldn't all be FIFO
- 
- Many/most of networking/systems research topics are related to QoS
    - latency, bandwidth, drop, differentiation
    - constraints / objectives
    - “local” or system-wide perspectives
  - Also working on *caching* problems
    - what to cache, when, size/value awareness, cost models, ...
    - local / distributed, self-adjusting
  - You're all welcome to stop by and chat anytime
    - also about research... ☺

# Next Semester (Undergraduate Course)



Thank You!

`sgabriel@bgu.ac.il`