The background of the slide features a light gray globe centered behind the text. Overlaid on the globe and the entire slide is a pattern of green binary code (0s and 1s) arranged in a perspective grid that recedes towards the top right.

# (Some Thoughts on) Cloud Security

Gabriel Scalosub

# Outline

- Introduction
  - Concerns and Threats
- Approaches to Securing the Cloud
- Concrete Example: Cache Side-Channel Attacks
  - Framework and General Methods
  - Some Variants

# Security: (The) Major Concern in Cloud

- Multi-tenancy vs. isolation
  - It's all about the money, remember?
- Resource sharing
  - Physical resources
  - SW libraries (kernel host)
  - Networks
- VMs vs. containers:
  - Stronger isolation (VMs)
  - More agility (containers)
  - Attack surface: large (VMs) vs small(er) (containers)
  - A lot of “philosophical” debate...





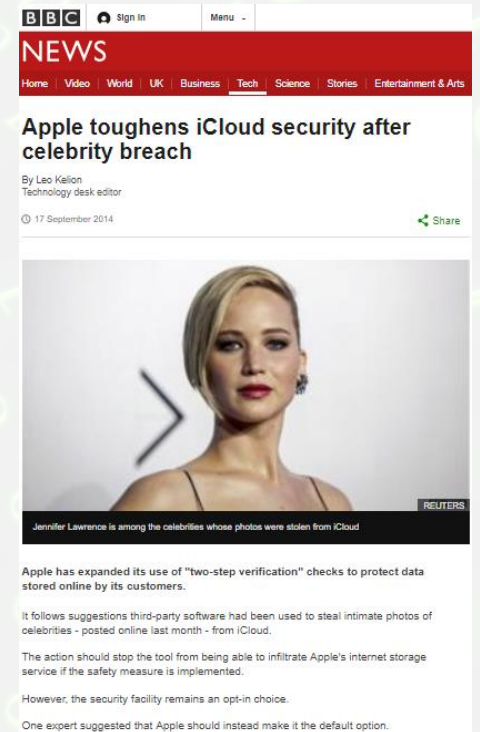
# Security: (The) Major Concern in Cloud

- Data breaches
- Insecure APIs
  - E.g., Web Applications
    - Cross-side scripting, SQL injection, ...

*[http://vulnerable-website.com/search?search\\_term='<script>\(bad things happen here\)</script>'](http://vulnerable-website.com/search?search_term='<script>(bad things happen here)</script>')*

- Distributed denial-of-service (DDoS)
- Side channel attacks
  - Stay tuned...
- Many more

- Is that it?
- What about trust?
  - Can I trust my cloud provider with my data?



Source: [www.bbc.com](http://www.bbc.com)

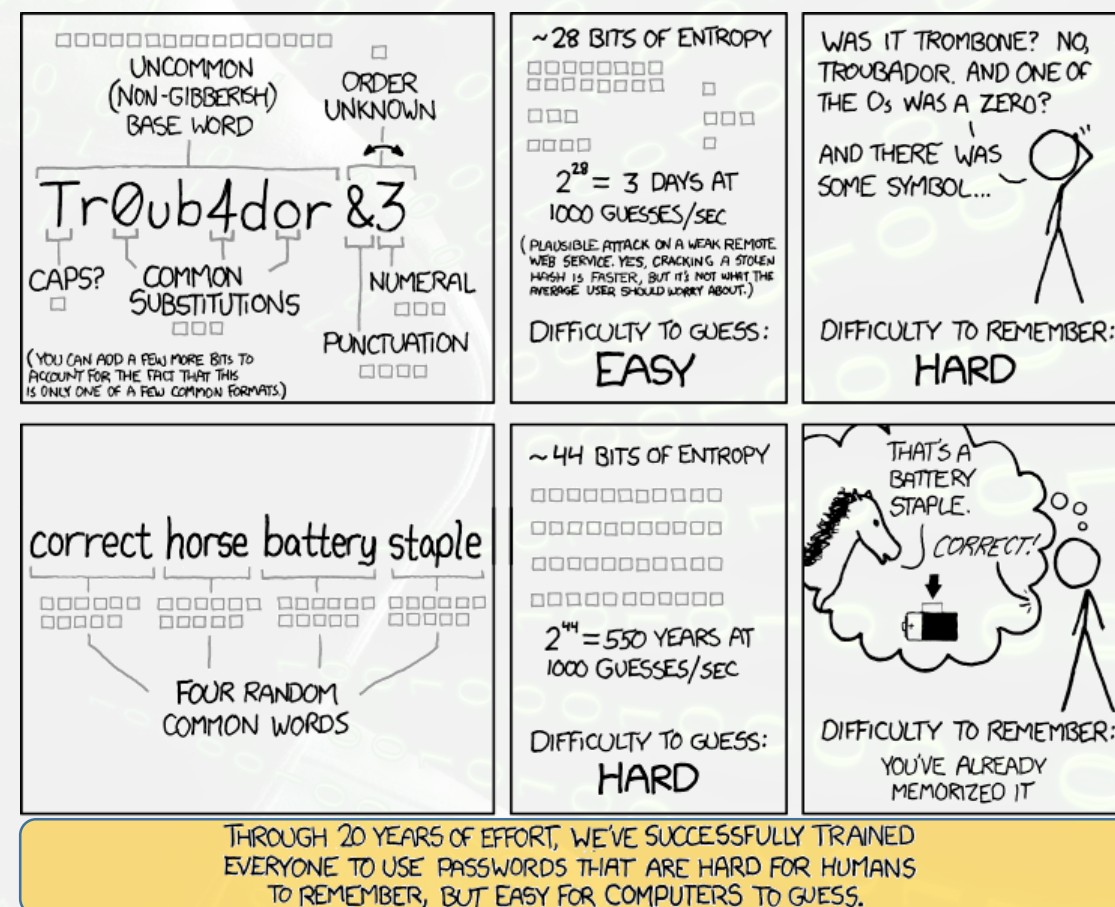
# Outline

- Introduction
  - Concerns and Threats
- Approaches to Securing the Cloud
- Concrete Example: Cache Side-Channel Attacks
  - Framework and General Methods
  - Some Variants



# Data-centric Security

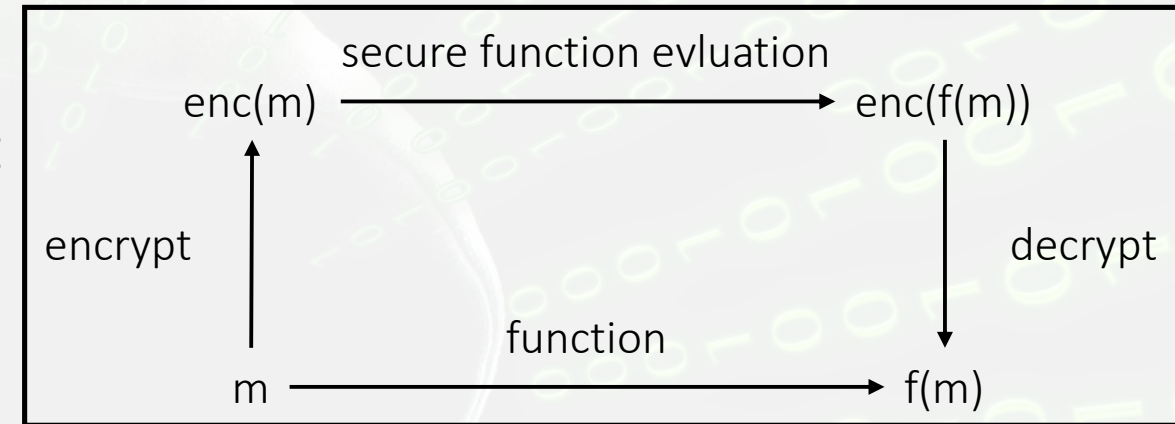
- Application-level
  - Authentication
    - 2-phase (a password is not enough)
  - Identity management & credentials
  - Data-loss
    - Backups...
  - Risk management
    - Supply chain & 3<sup>rd</sup> party
    - Know thy-own vulnerabilities
  - Usually based on cryptography
- Also system-level
  - E.g., secret sharing (information theoretic)



Source: <https://www.explainxkcd.com/>

# Secured Cloud Data: Homomorphic Encryption

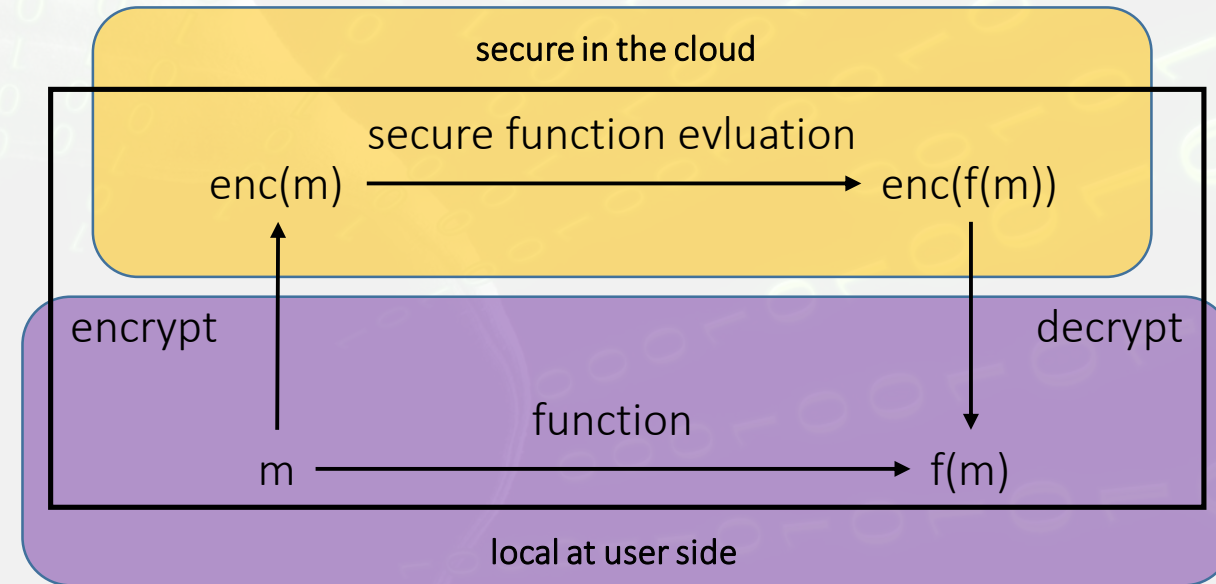
- Allows computation on ciphertext
- Simple example (not really secure...):
  - Encryption by  $E(x) = 2x$
  - Decryption by  $D(x) = x/2$
  - Plaintext inputs  $x_1 = 3$  and  $x_2 = 4$ 
    - Encrypted ciphertext data:  $E(x_1) = 6, E(x_2) = 8$
    - Operation on ciphertext:  $2 * E(x_1) + 3 * E(x_2) = 36$
    - Decrypt operation on ciphertext:  $D(2 * E(x_1) + 3 * E(x_2)) = 36/2 = 18 = 2 * x_1 + 3 * x_2$
- Real example:
  - Private patients data encrypted in a DB in the cloud
  - Pharmaceuticals want to compute statistics on patients data (e.g., AVG, distribution,...)
  - Draw conclusions from statistics, without access to plaintext data, only to ciphertext





# Secured Cloud Data: Homomorphic Encryption

- Scheme requirement:
  - Secure!! (crypto / information theoretic)
- Partial scheme:
  - Supports just one operation (e.g., +)
- Full scheme:
  - Supports two operations (e.g., +, ·)
  - Enables computing any function
  - Current SOTA: very slow compared to plaintext operation
    - More than  $\times 10^3$  slower...
- Applications:
  - Secure voting
  - Private information retrieval (PIR)
    - Query a DB without anyone knowing what your *query* was
  - Machine learning on anonymized data





# Infrastructure-centric Security

- Firewalls
- Intrusion detection systems (IDS)
- Antivirus (AVs)
- (No)-deduplication
  - Duplicate shared libraries -> deduplicate (more efficient) -> no de-deuplication
  - Increase isolation
    - Less efficient
- Secure against VM escape
  - VM is able to access/attack the host-OS / hypervisor
  - Obtain privileged access
    - Hypervisor / co-located VM

# Outline

- Introduction
  - Concerns and Threats
- Approaches to Securing the Cloud
- Concrete Example: Cache Side-Channel Attacks
  - Framework and General Methods
  - Some Variants

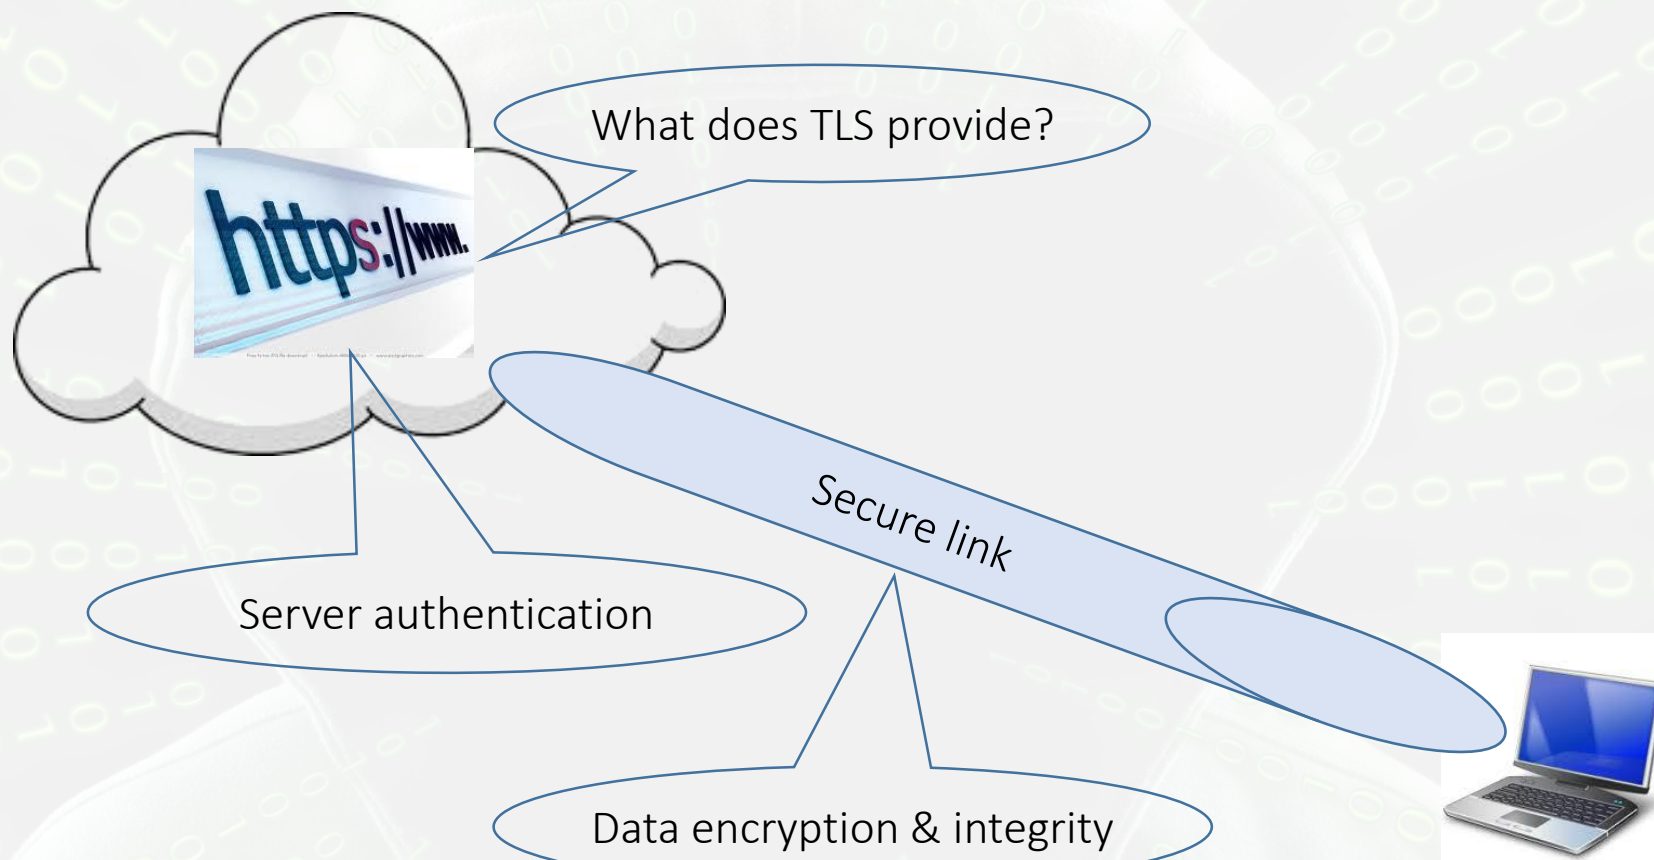
# Transport Layer Security (TLS) / HTTPS



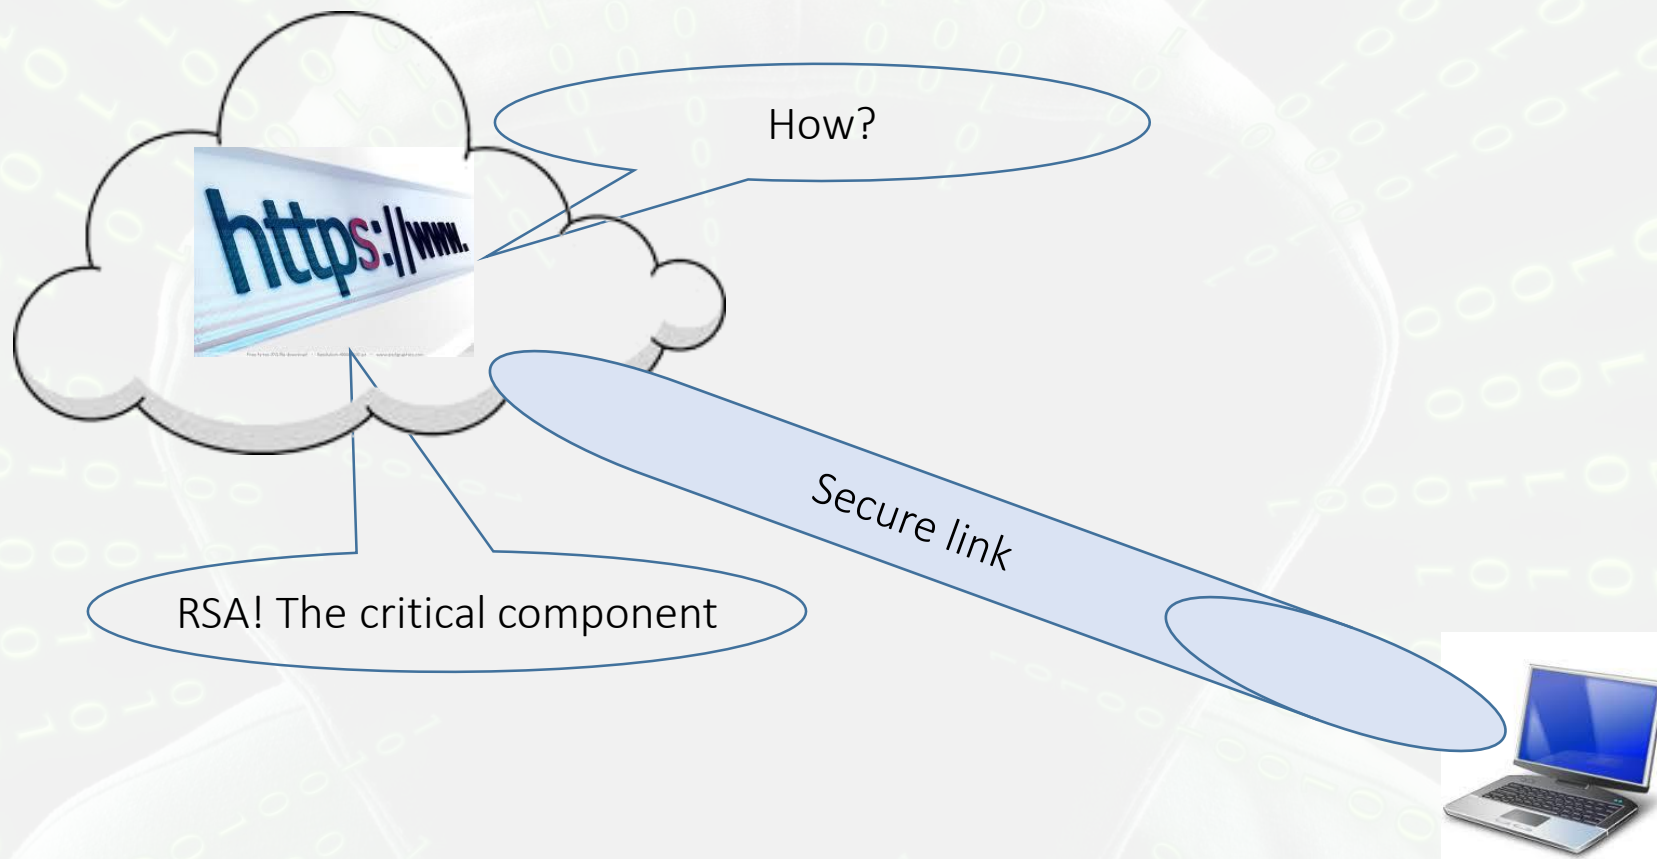
Here and there...



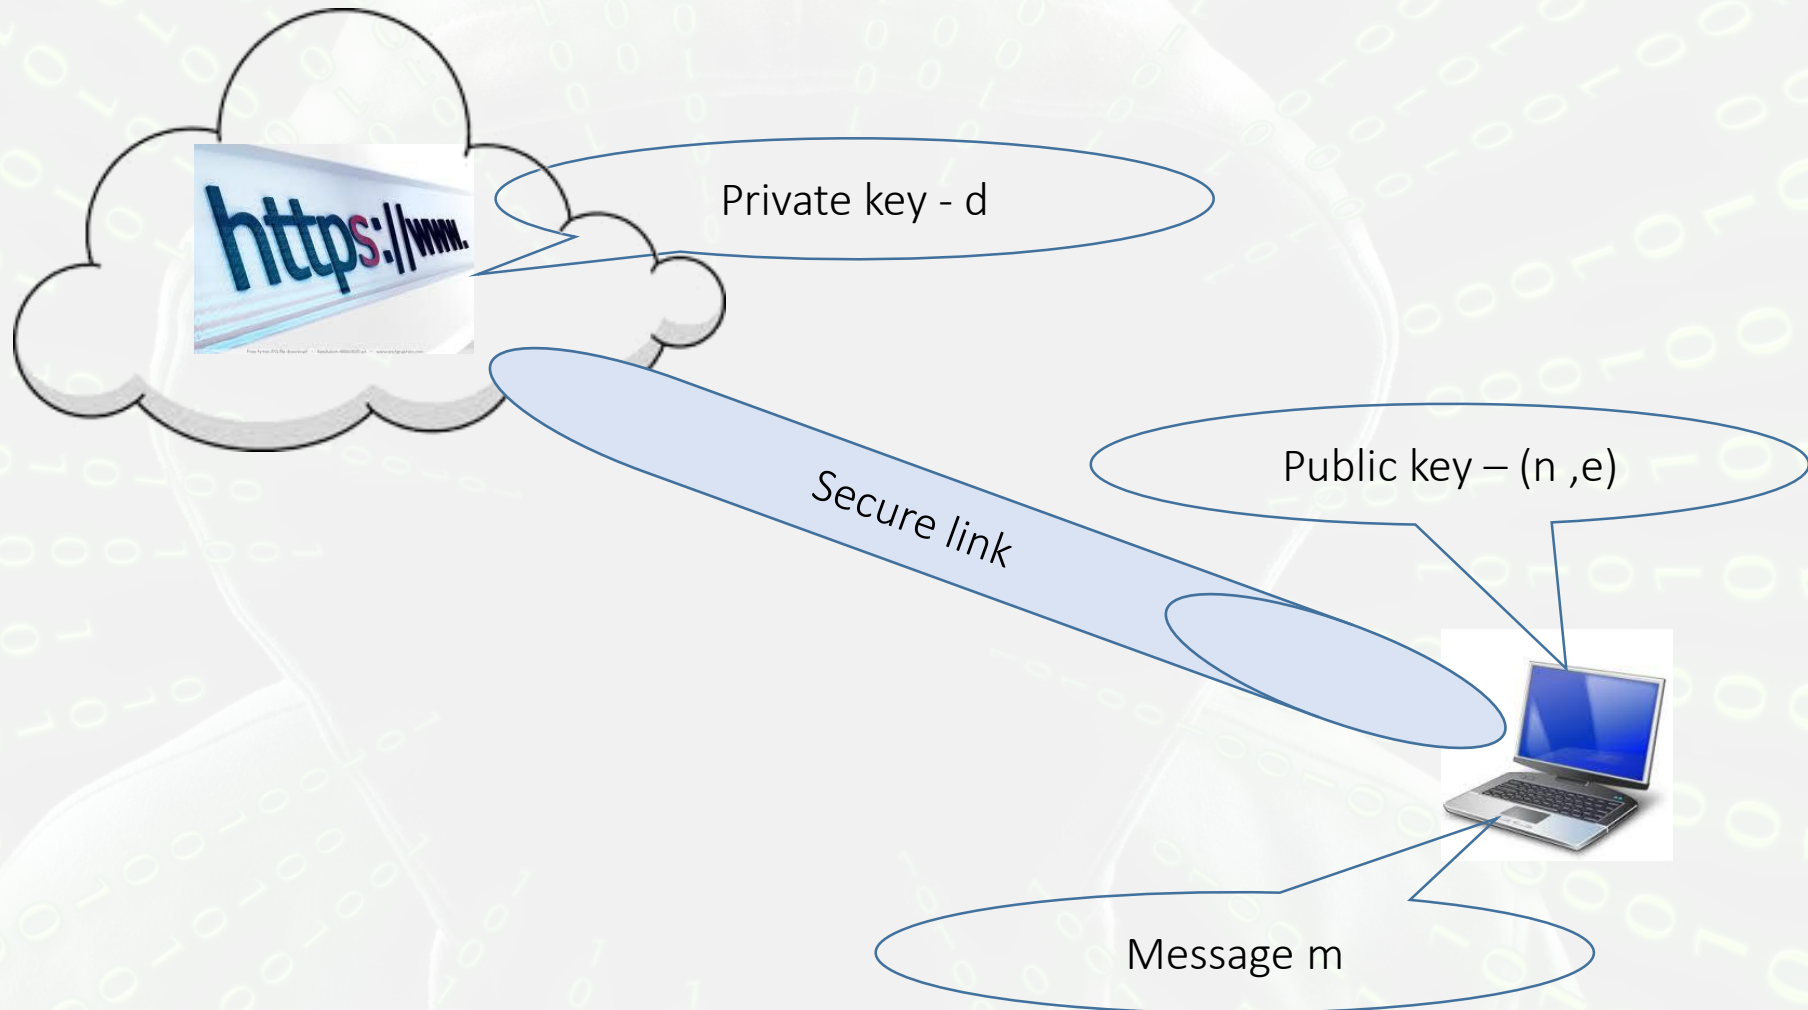
# Transport Layer Security (TLS) / HTTPS



# Transport Layer Security (TLS) / HTTPS

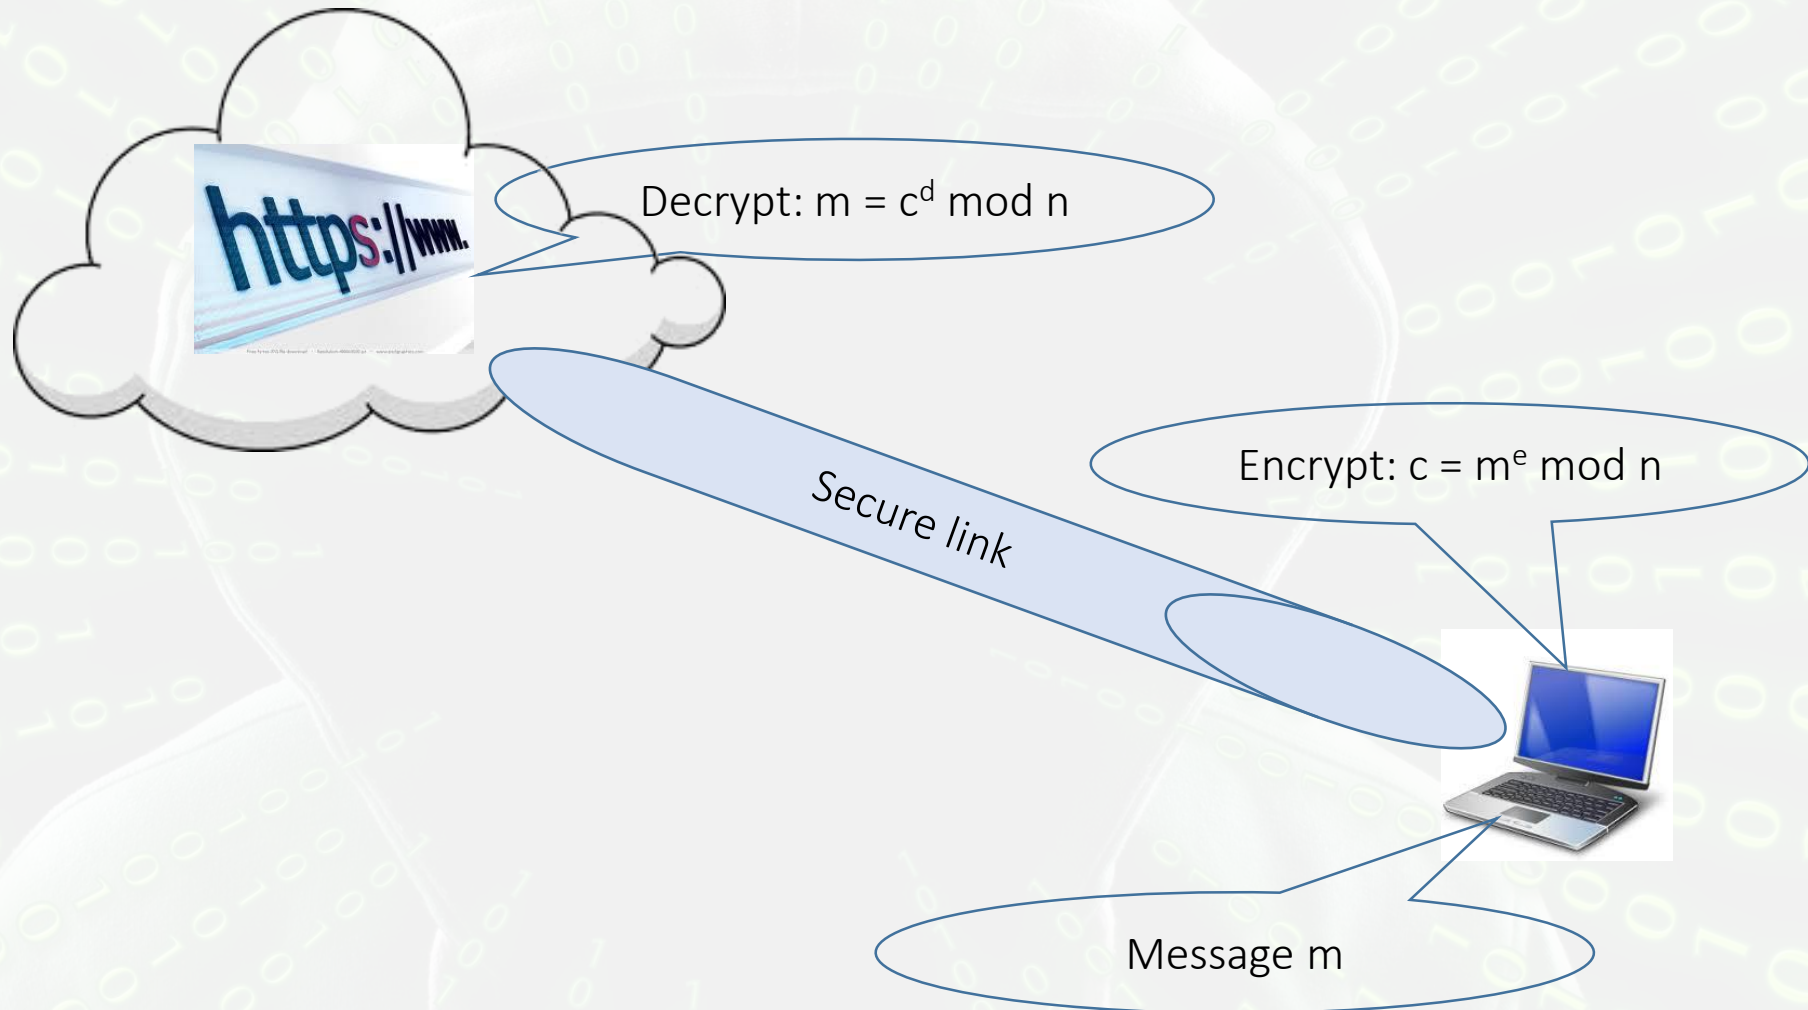


# TLS / HTTPS (Specifically, RSA)

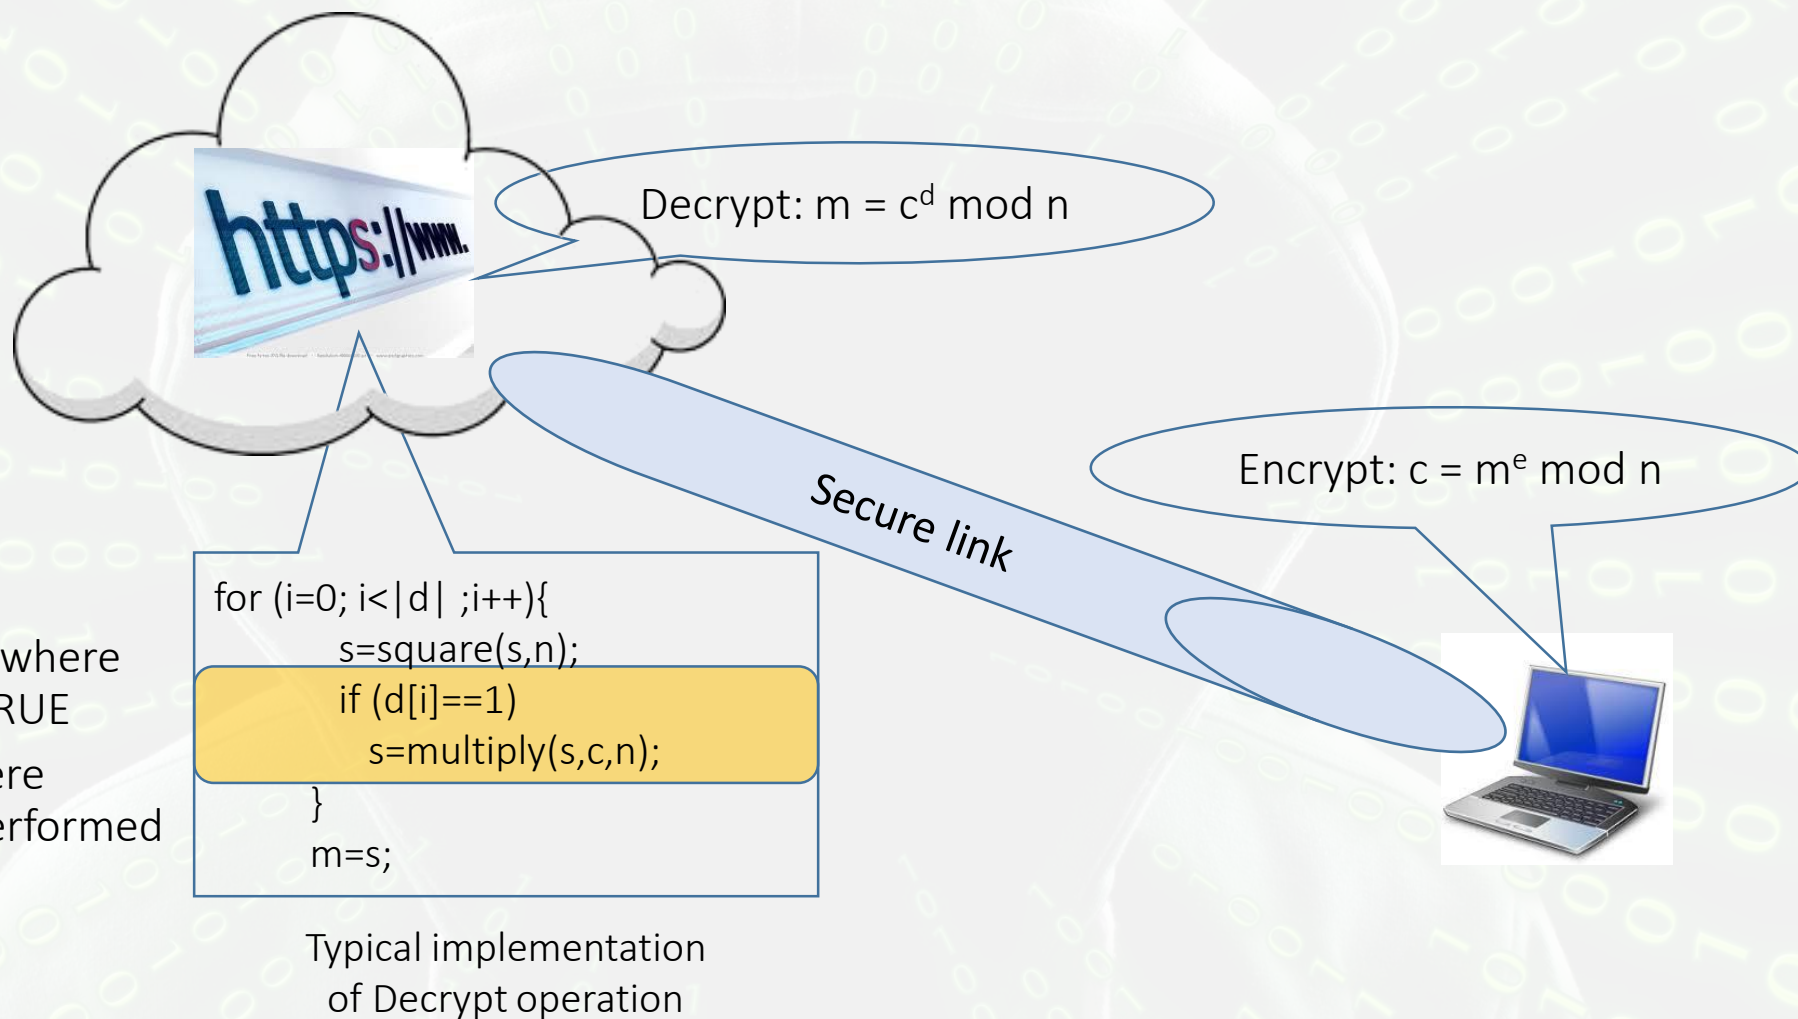




# TLS / HTTPS (Specifically, RSA)

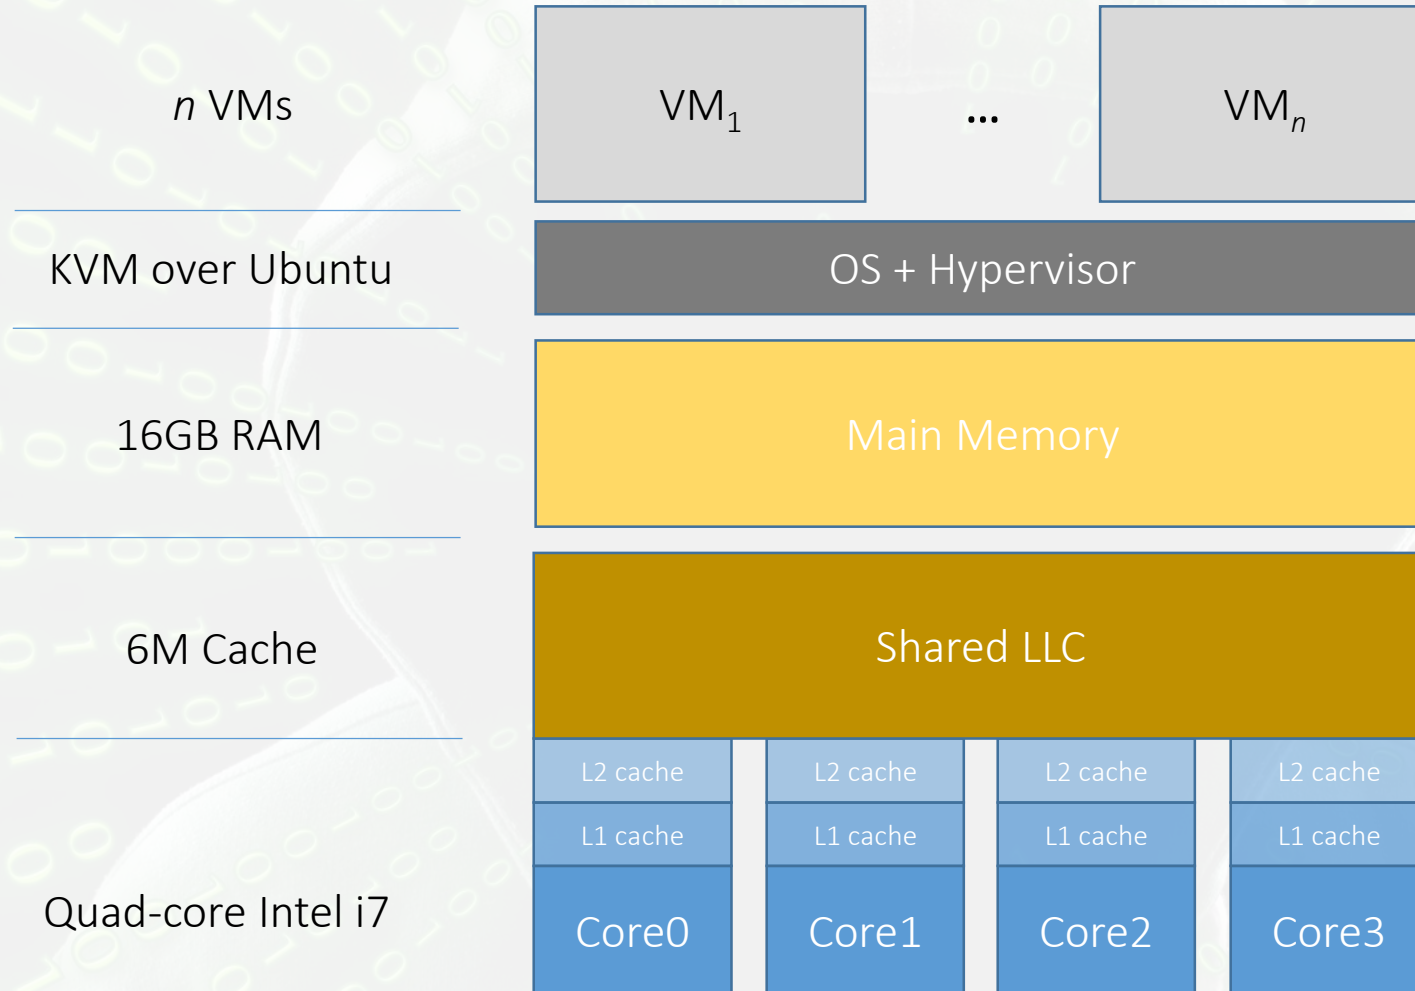


# TLS / HTTPS (Specifically, RSA)



- Goal:
  - Identify d-bits where statement is TRUE
  - I.e., d-bits where multiply() is performed

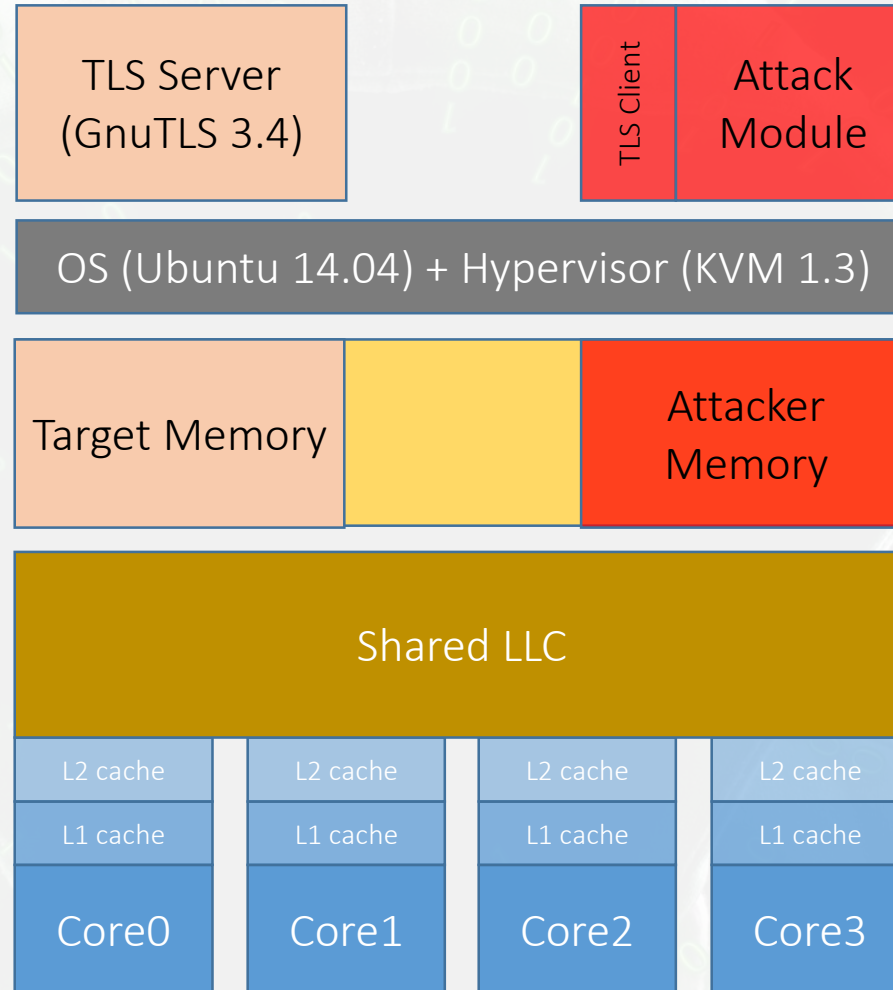
# A Concrete System



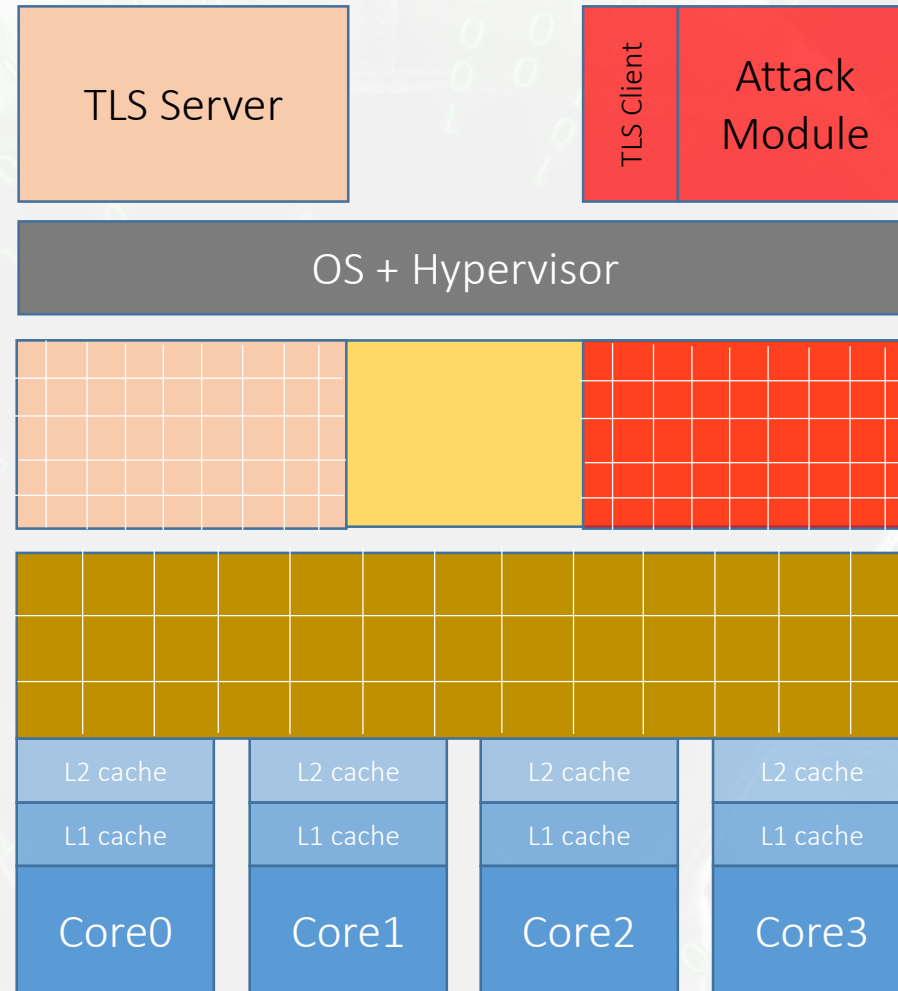


# A Concrete System

Secret key:  
 $|d|=4096$  bit



# A Concrete System



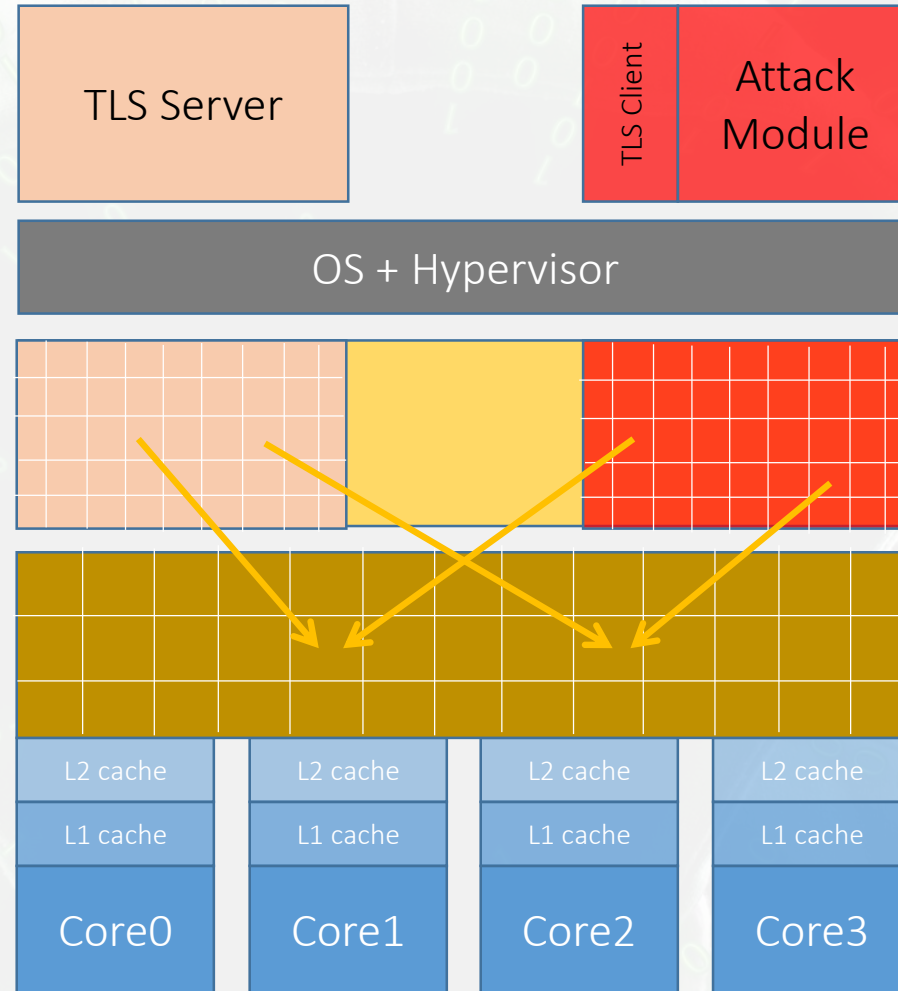
6M Cache = 8192 sets

12-way associative

- 12 lines / set

# A Concrete System: Memory vs. Cache

Many-to-one





# A Concrete System: Memory vs. Cache

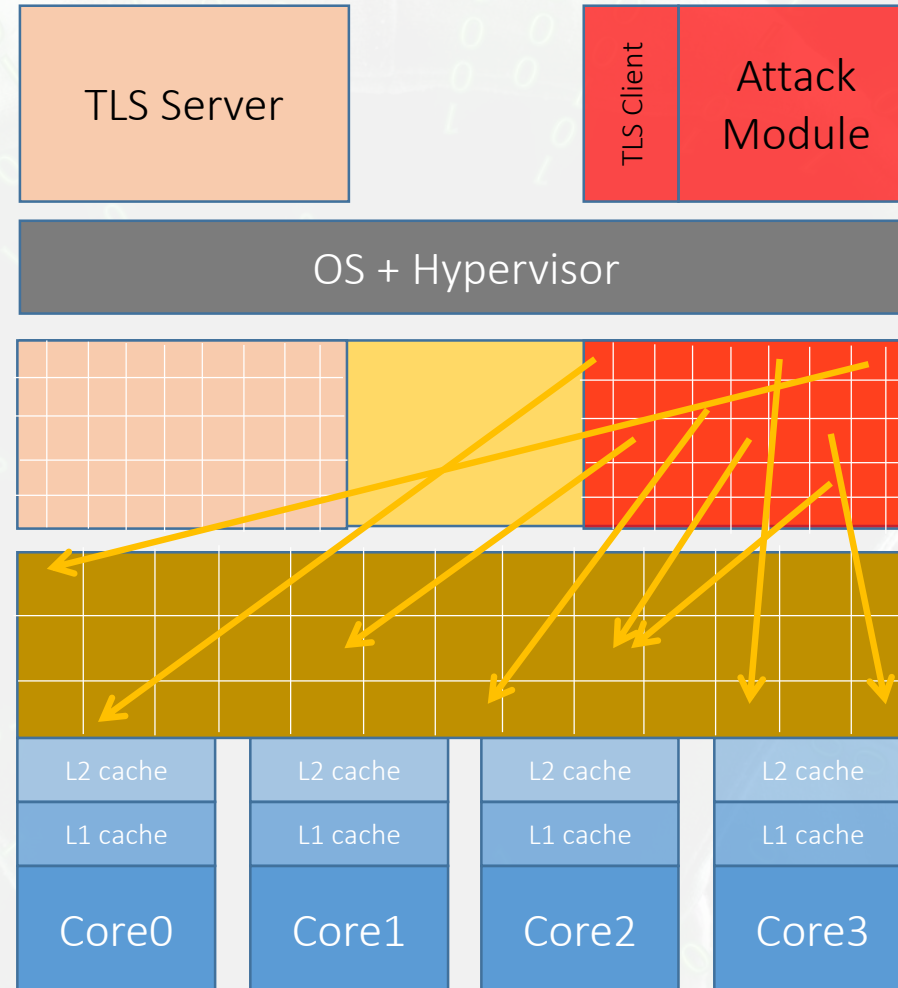
Attacker knows  
its own mapping

Builds collection of  
12 lines / set

Enables the attacker to  
“Flush” entire set

Prime&Probe:

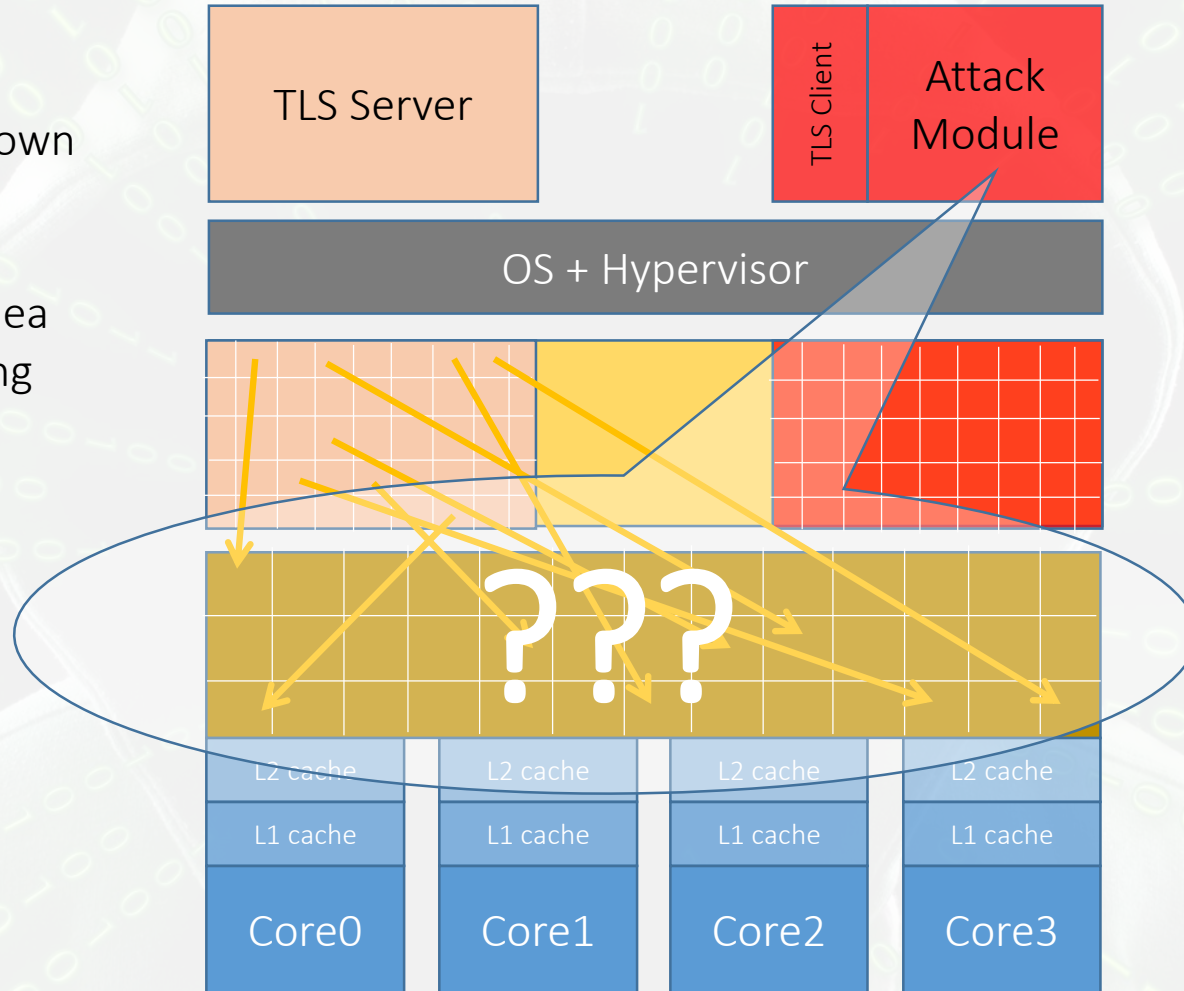
- Prime: “Flush”
- Wait (for target activity)
- Probe: check for miss



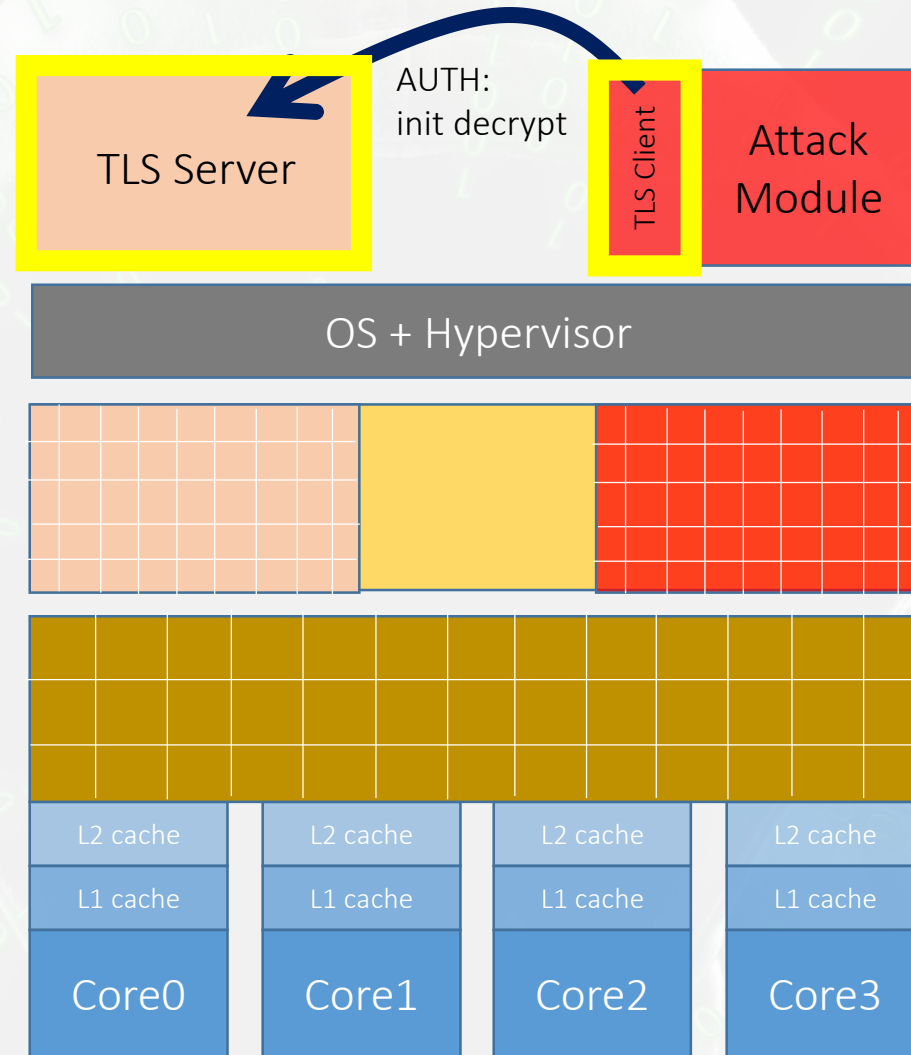
# A Concrete System: Memory vs. Cache

Target also has it own mapping

Attacker has no idea of target's mapping

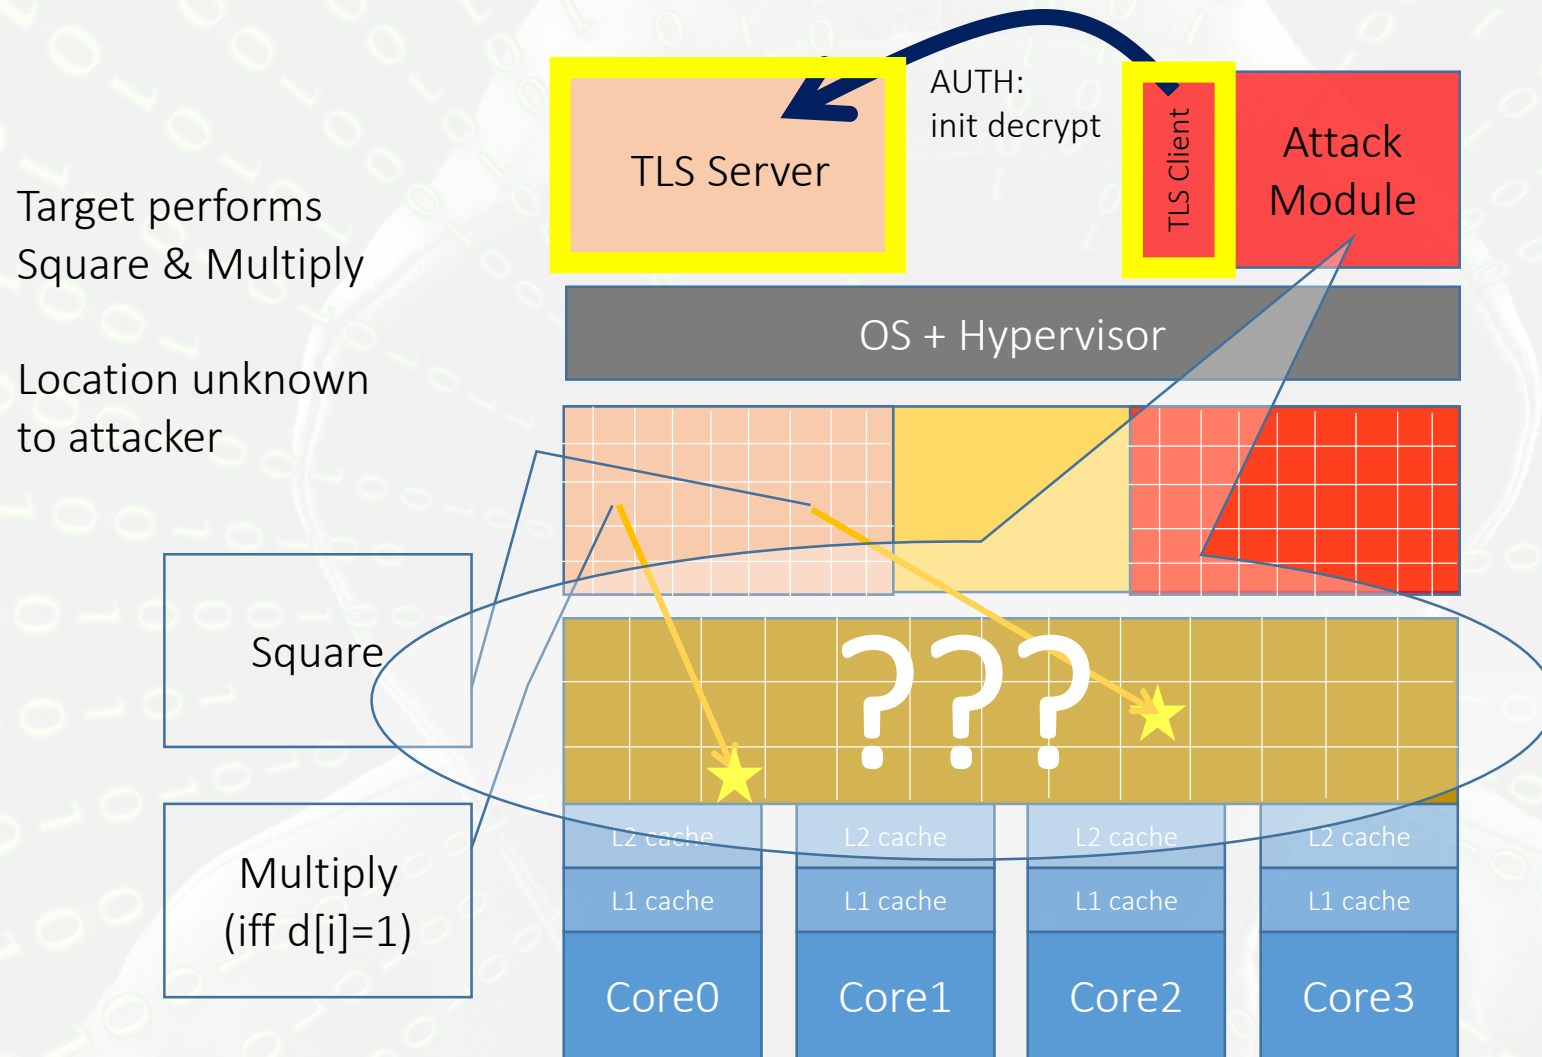


# Attack Initialization





# Attack: Under the Hood



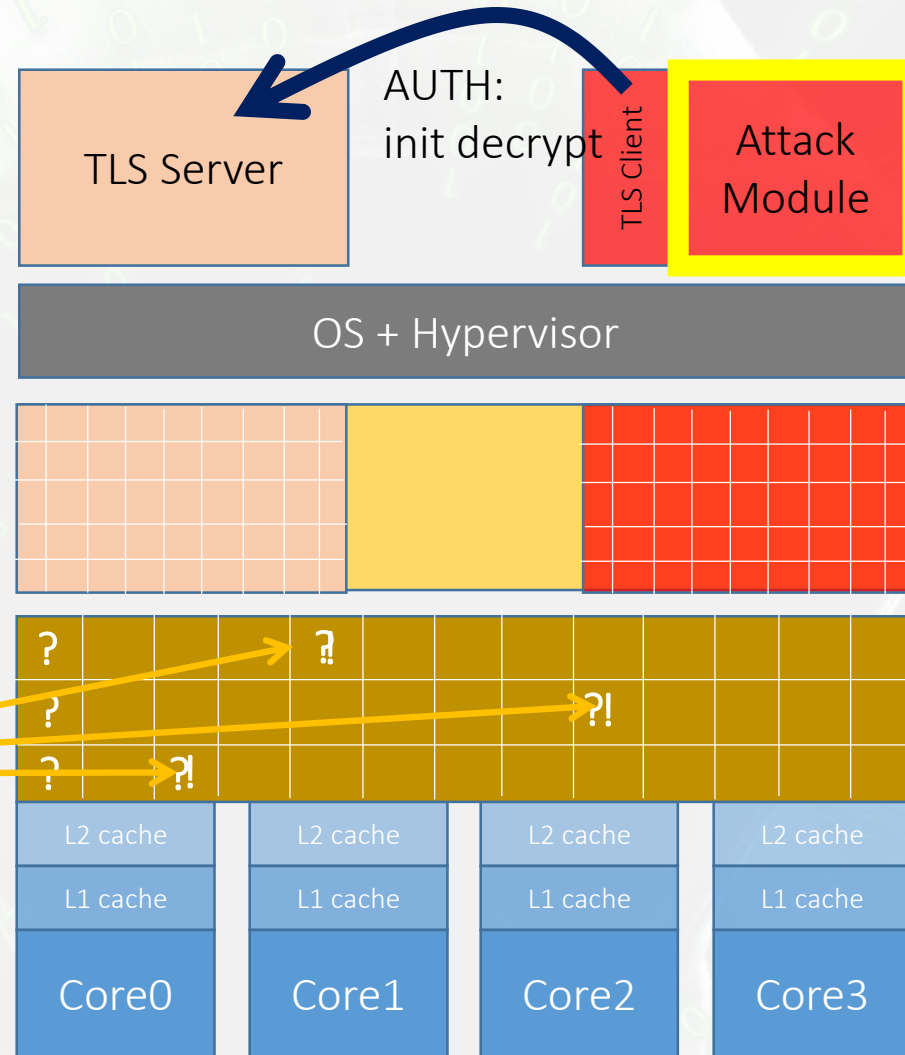
# Attack: Prime & Probe

Attack Module:

Scans cache sets  
for activity

Prime & Probe  
Cache hit/miss

Identify active sets  
(by score)

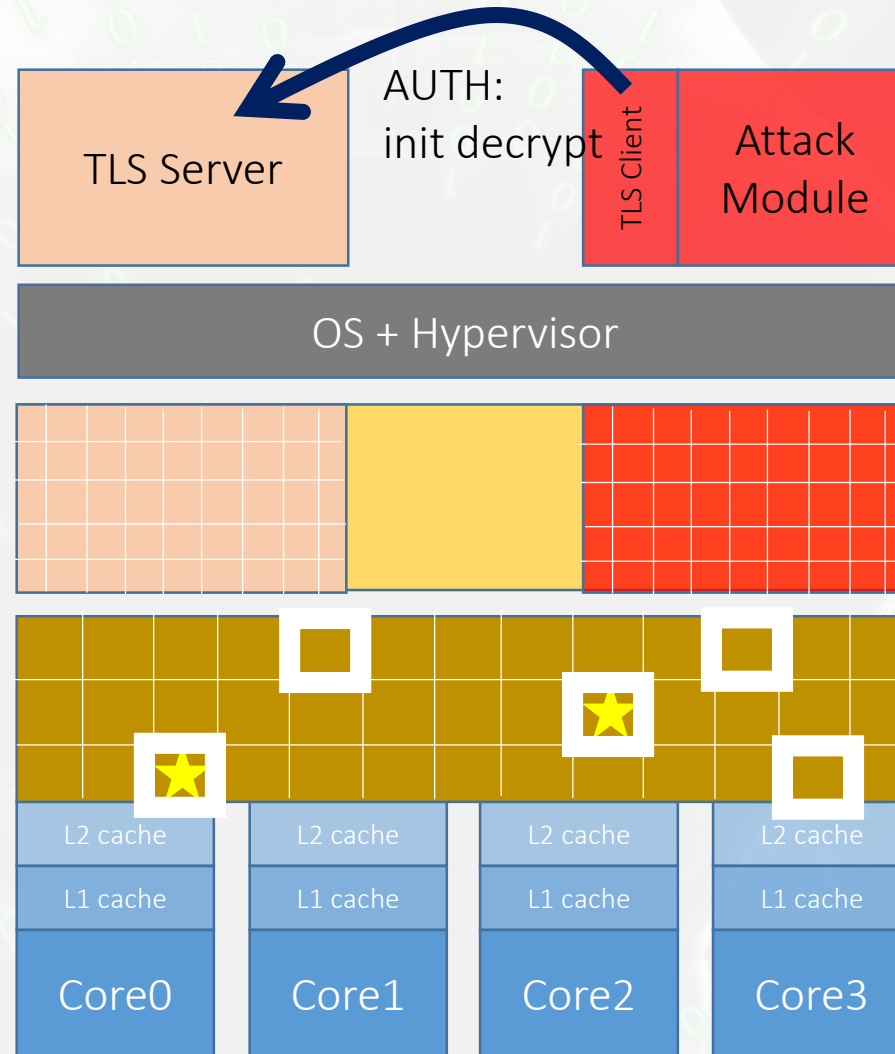


# Attack: Prime & Probe

List of active sets

Attacker doesn't know which are "correct"

Test each active set for signature traces



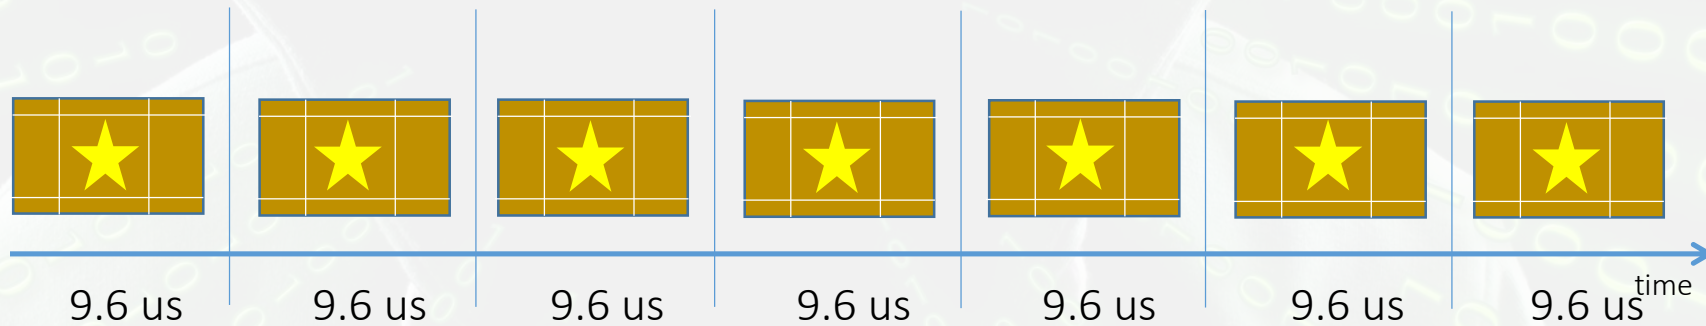


# Attack: Extract Samples

Suspected “Multiply” set

Used if-and-only-if secret key bit  $d[i]=1$

Target takes 9.6us to process a single d-bit

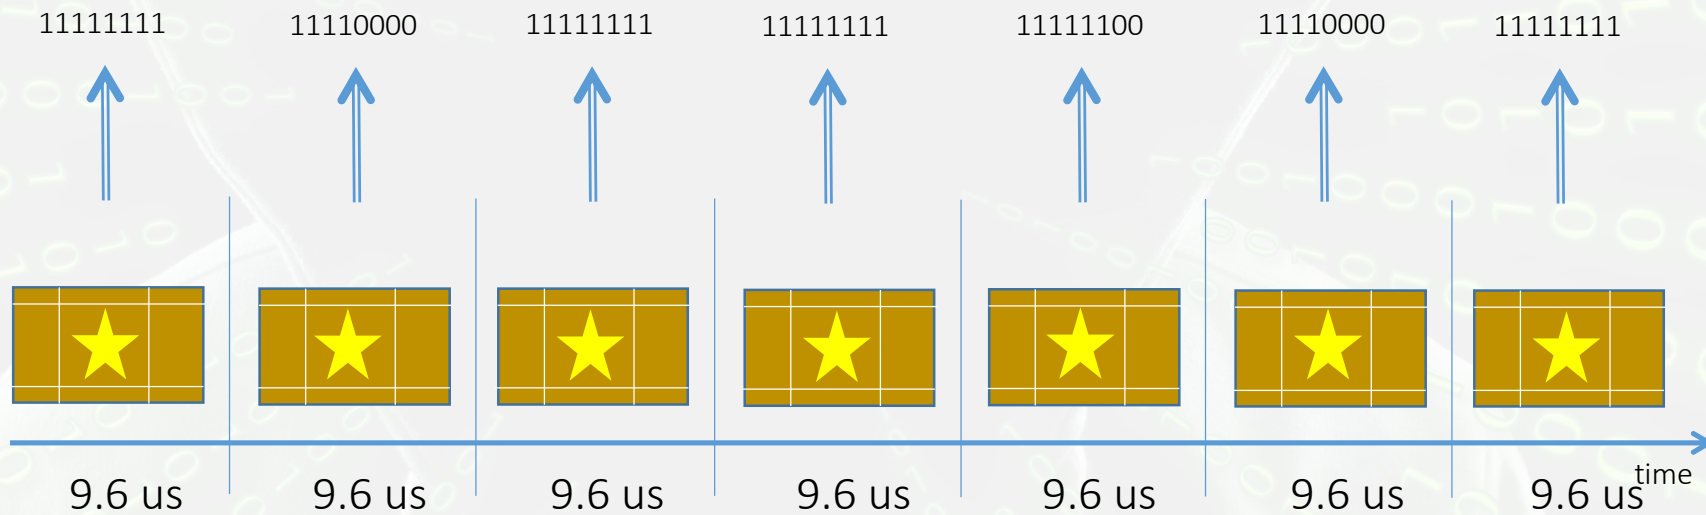


# Attack: Extract Samples

## Sample: Prime&Probe

Extract samples from each d-bit

8 samples per d-bit  
(one sample every 1.2us)

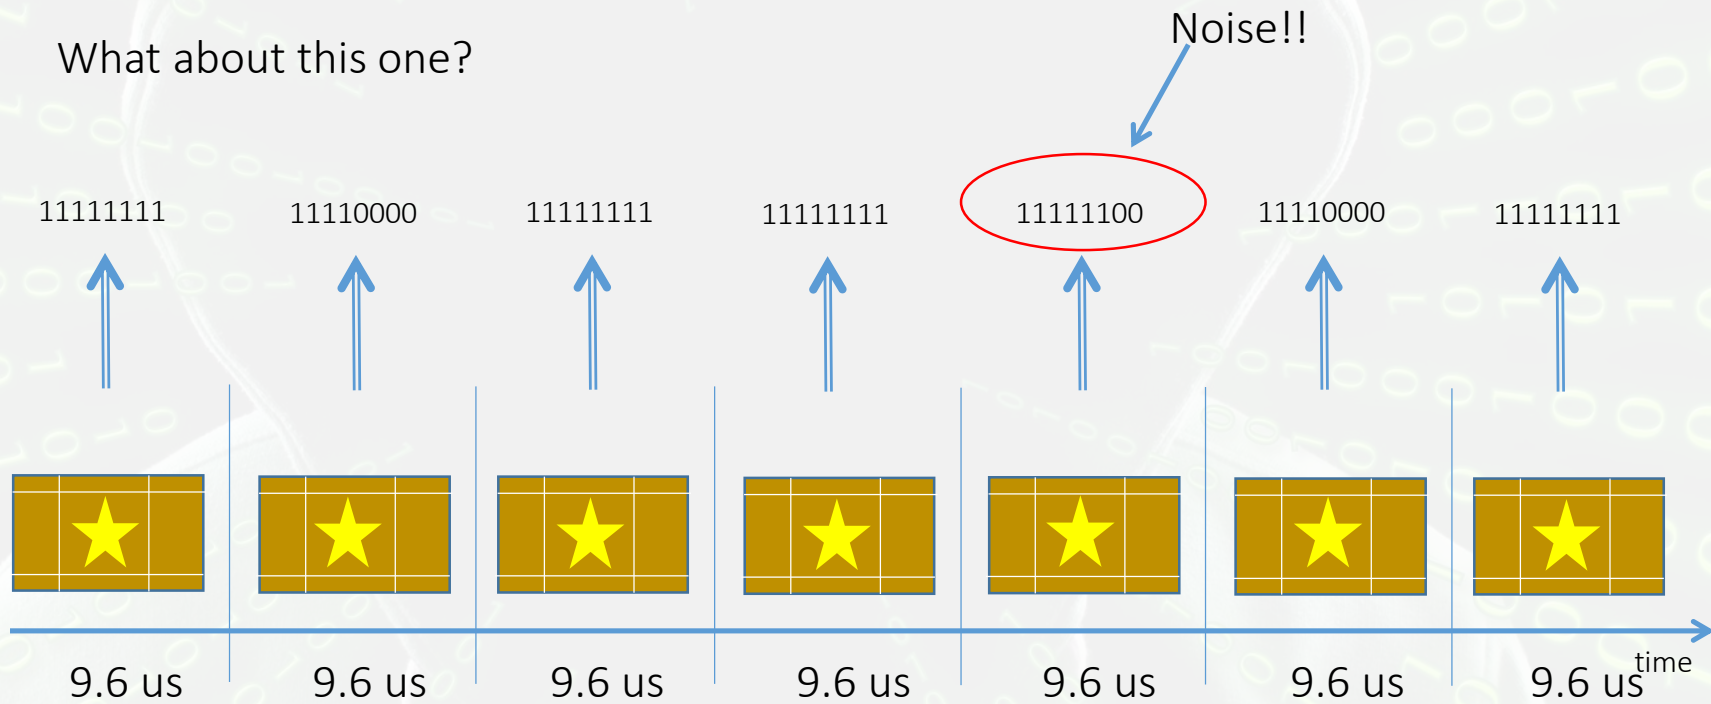


# Attack: Extract Samples

Good vs. bad patterns

Can you tell which are which?

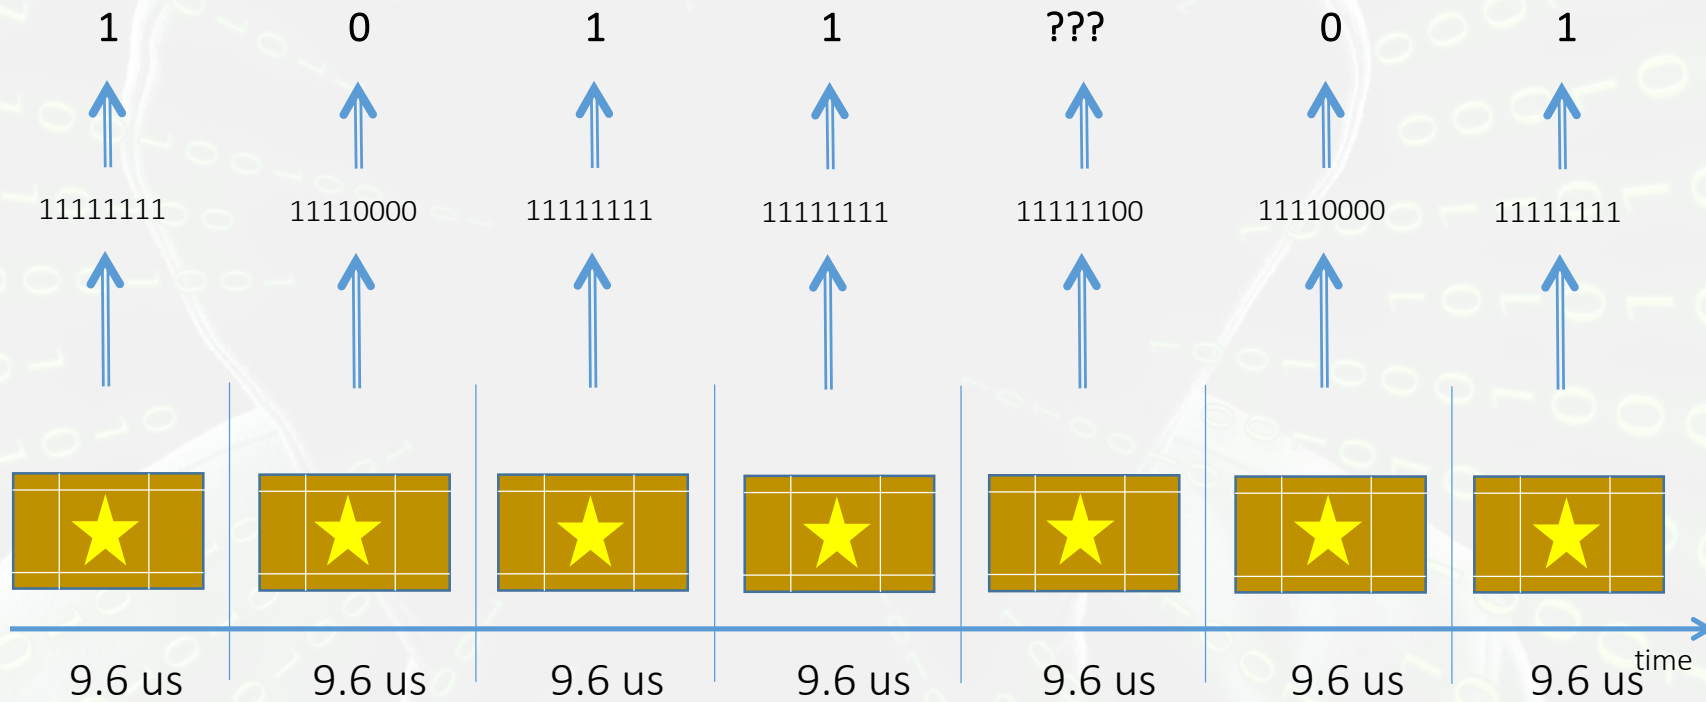
What about this one?





# Attack: From Samples to Bits

Translate samples to candidate d-bits



# Attack: From Bits to Complete Keys

“Good” candidate sequences:

- Long, few missing bits

```
1110100100010100100001010100101010100111111000101000111010100101101100
1000111010010001010010000101010010101010011011100010100011101010010110
0000111000111010010001010010110101010010101010011111100010100011101010
0000111000111010010001010010110101010010101010011111100010100011101010
0011101001000101011000010101001010101001111110001010001110101001011011
0010001110100010010100100001010100101010100111111000101000111010100101
1000111010010001010010000101010010101010011111100010100011101010010110
```

# Attack: From Bits to Complete Keys

Align

```
1110100100010100100001010100101010100111111000101000111010100101101100
1000111010010001010010000101010010101010011011100010100011101010010110
0000111000111010010001010010110101010010101010011111100010100011101010
0000111000111010010001010010110101010010101010011111100010100011101010
0011101001000101011000010101001010101001111110001010001110101001011011
0010001110100010010100100001010100101010100111111000101000111010100101
1000111010010001010010000101010010101010011111100010100011101010010110
```



# Attack: From Bits to Complete Keys

# Majority

[illegible]

# Attack: From Bits to Complete Keys

Candidate key extracted

Validate?

- encrypt-decrypt...

```
1110100100010100100001010100101010100111111000101000111010100101101100
1000111010010001010010000101010010101010011111100010100011101010010110
00001110001110100100010100100001010100101010011111100010100011101010
00001110001110100100010100100001010100101010011111100010100011101010
0011101001000101001000010101001010101001111110001010001110101001011011
0010001110100100010100100001010100101010100111111000101000111010100101
1000111010010001010010000101010010101010011111100010100011101010010110
```

# Attack: From Bits to Complete Keys

After a few iterations...

“Correct” set is tested

True private key obtained!!

Target Key is:

0xdb3e5914a0b076f8288edf676

115b5c881dbebc87d6d5ce79a55

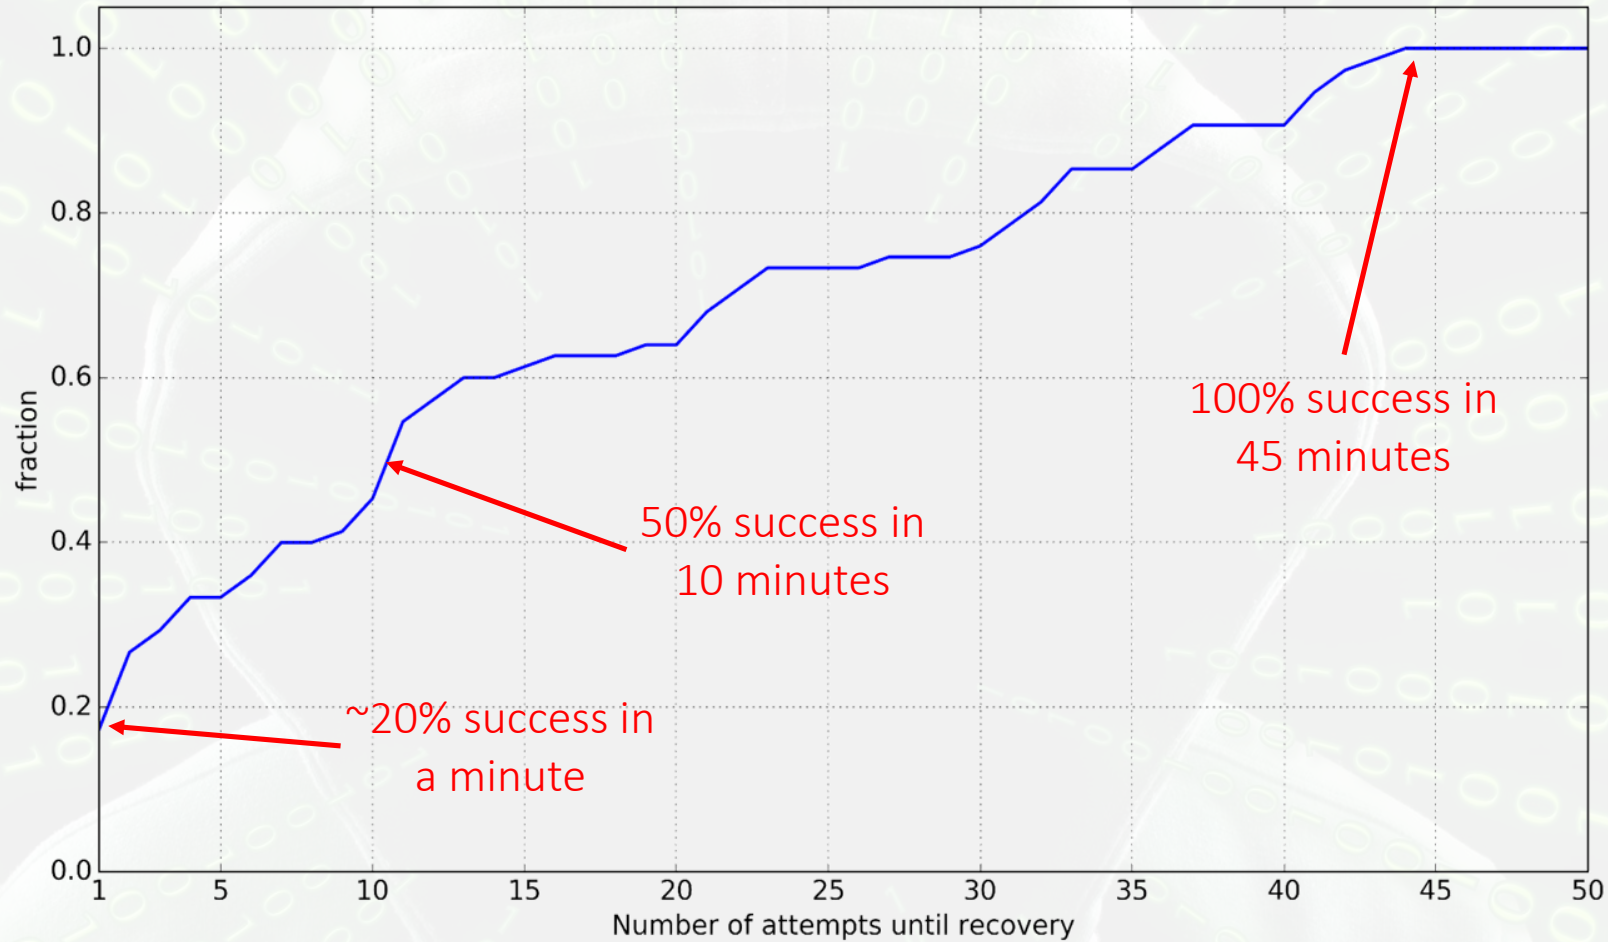
. . . . .

6f1bb6acd3b4d4454bbbbbd9752c

9f6ccf209054b3f803d7c63df11



# Attack: Success Rate



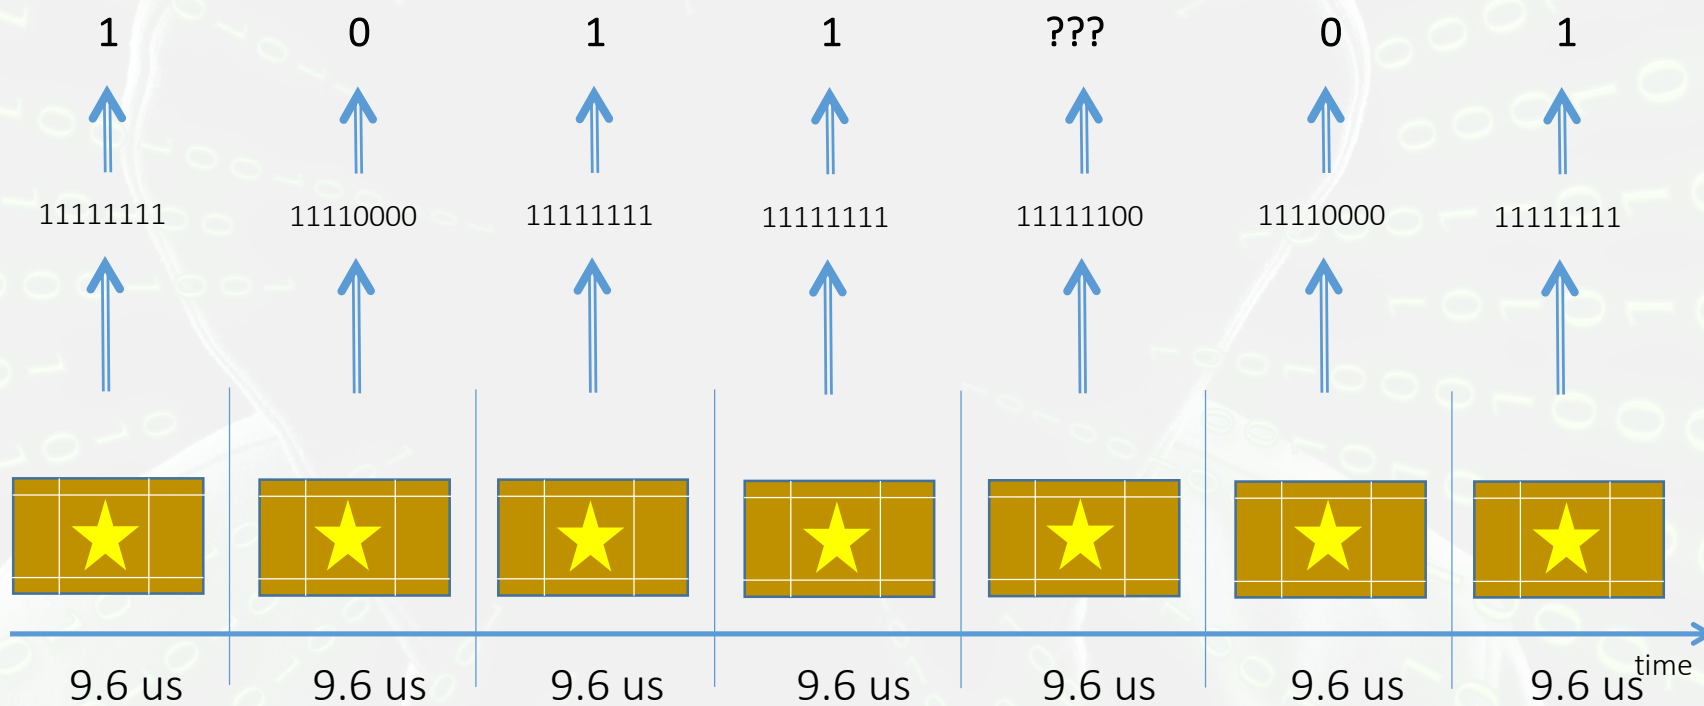
75 tests

50 attempts per test,  
200 samples per attempt

~1 minute per attempt

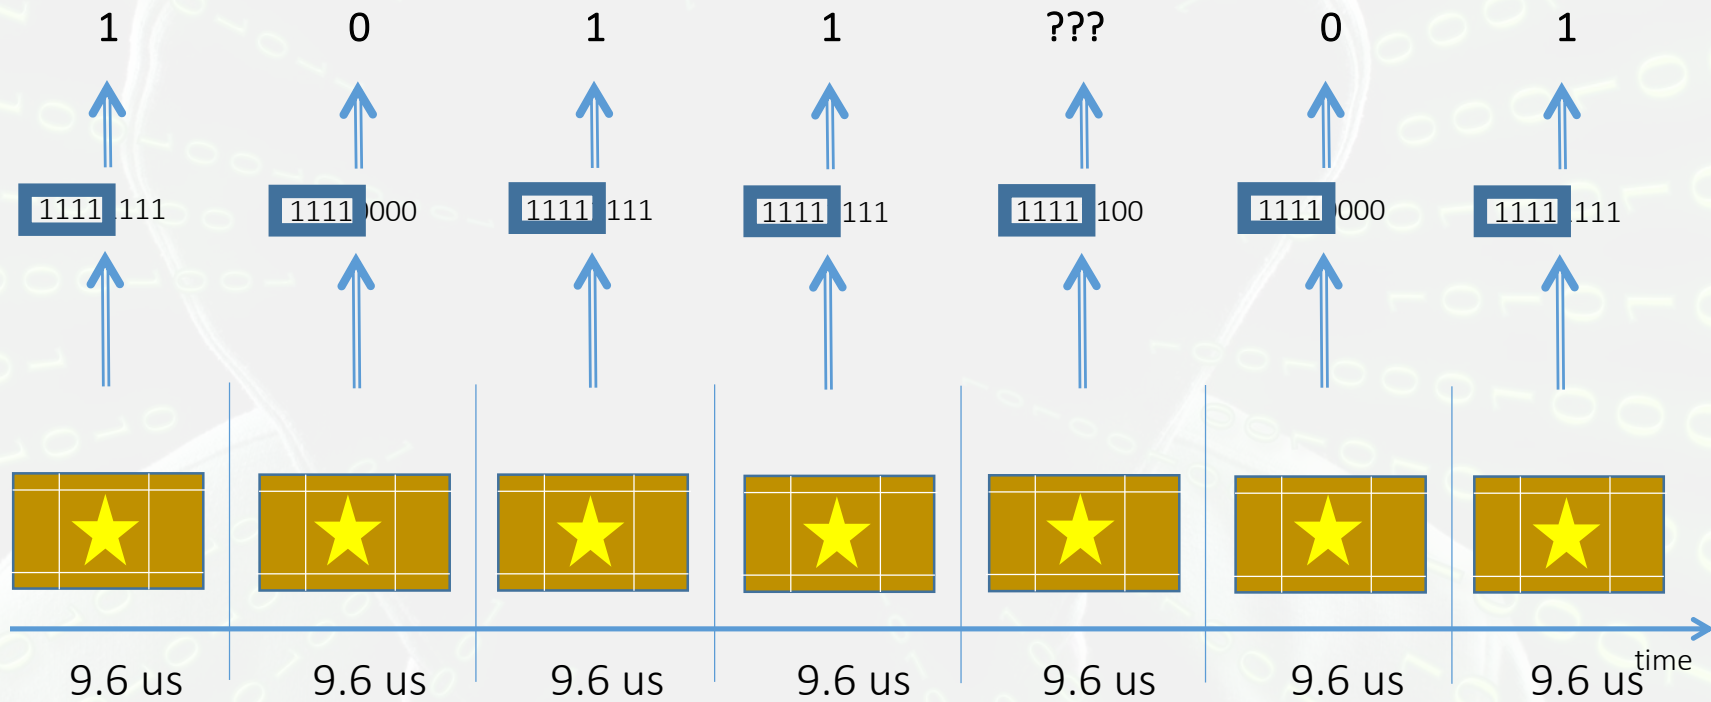
# Monitoring: How/What?

Recall samples gathered during the Prime&Probe attack



# Monitoring: How/What?

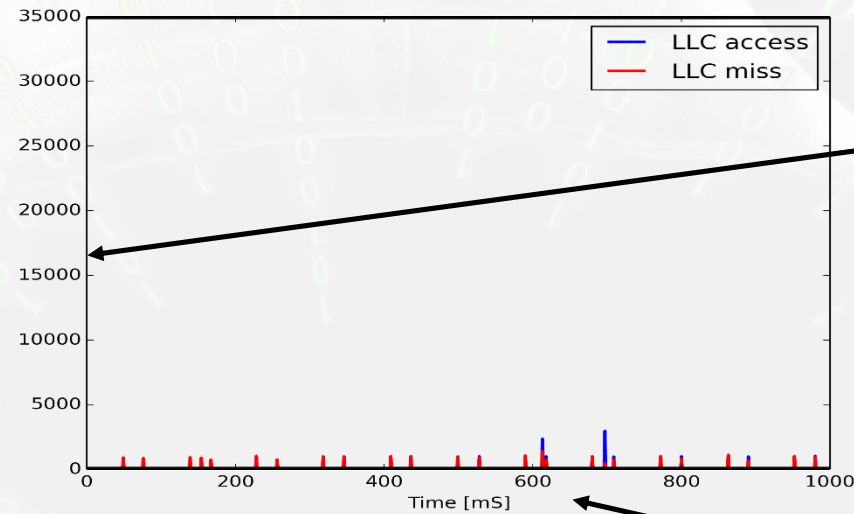
Target activity in correct set causes many misses





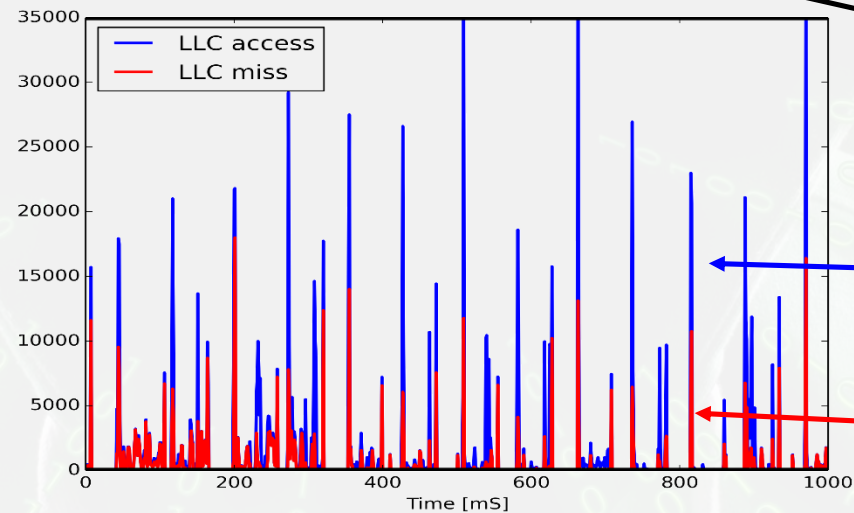
# Monitoring: Idle Attacker

Attacker Process



Y-axis:  
# LLC access/miss in  
time slot

Server Process

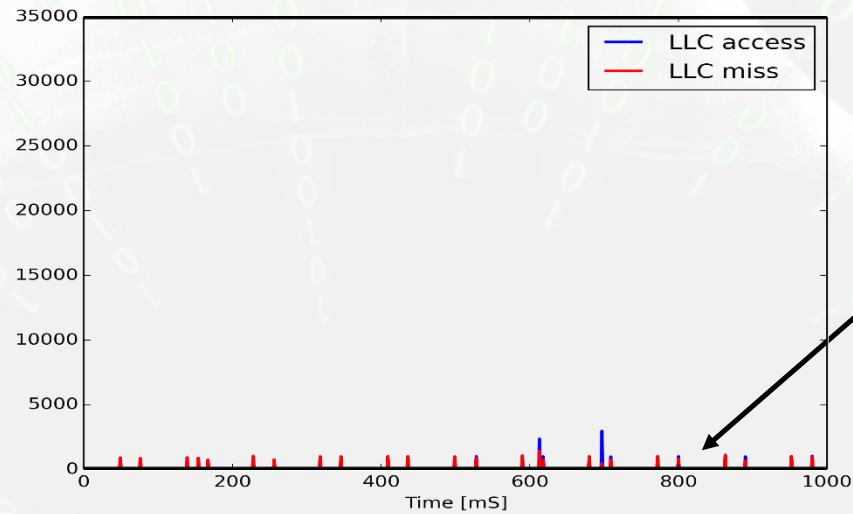


X-axis:  
time in ms

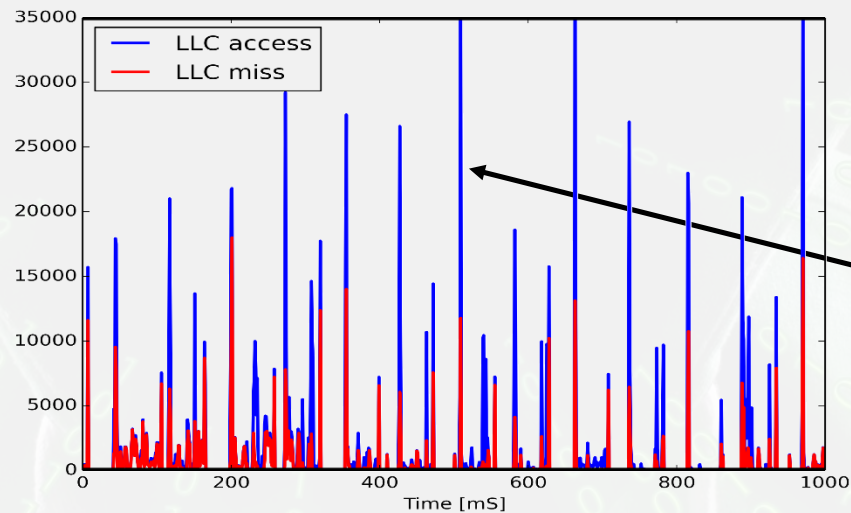
Blue line:  
cache access

Red line:  
cache miss

# Monitoring: Idle Attacker



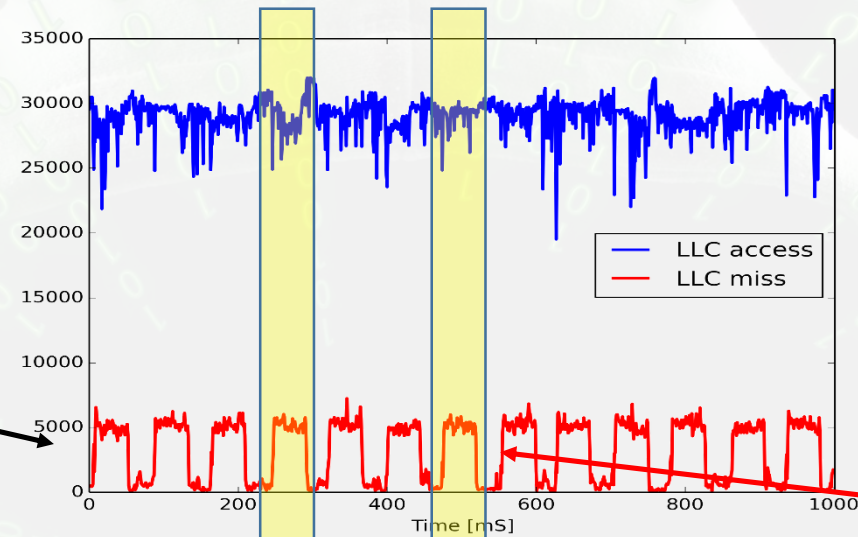
Attacker inactive –  
almost no cache  
activity



Server activity –  
bursts of LLC  
accesses & misses

# Monitoring: Active Attacker

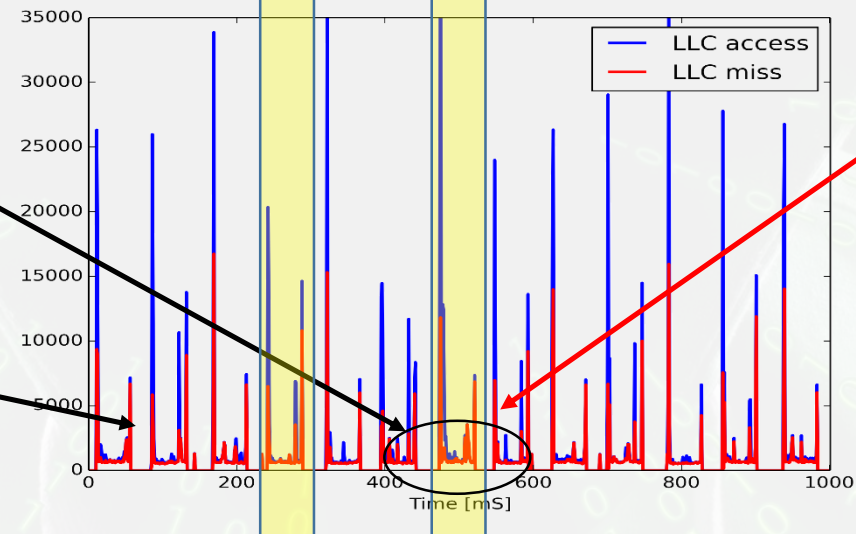
Nice...



server and attacker  
activity correlated!

RSA  
decryption

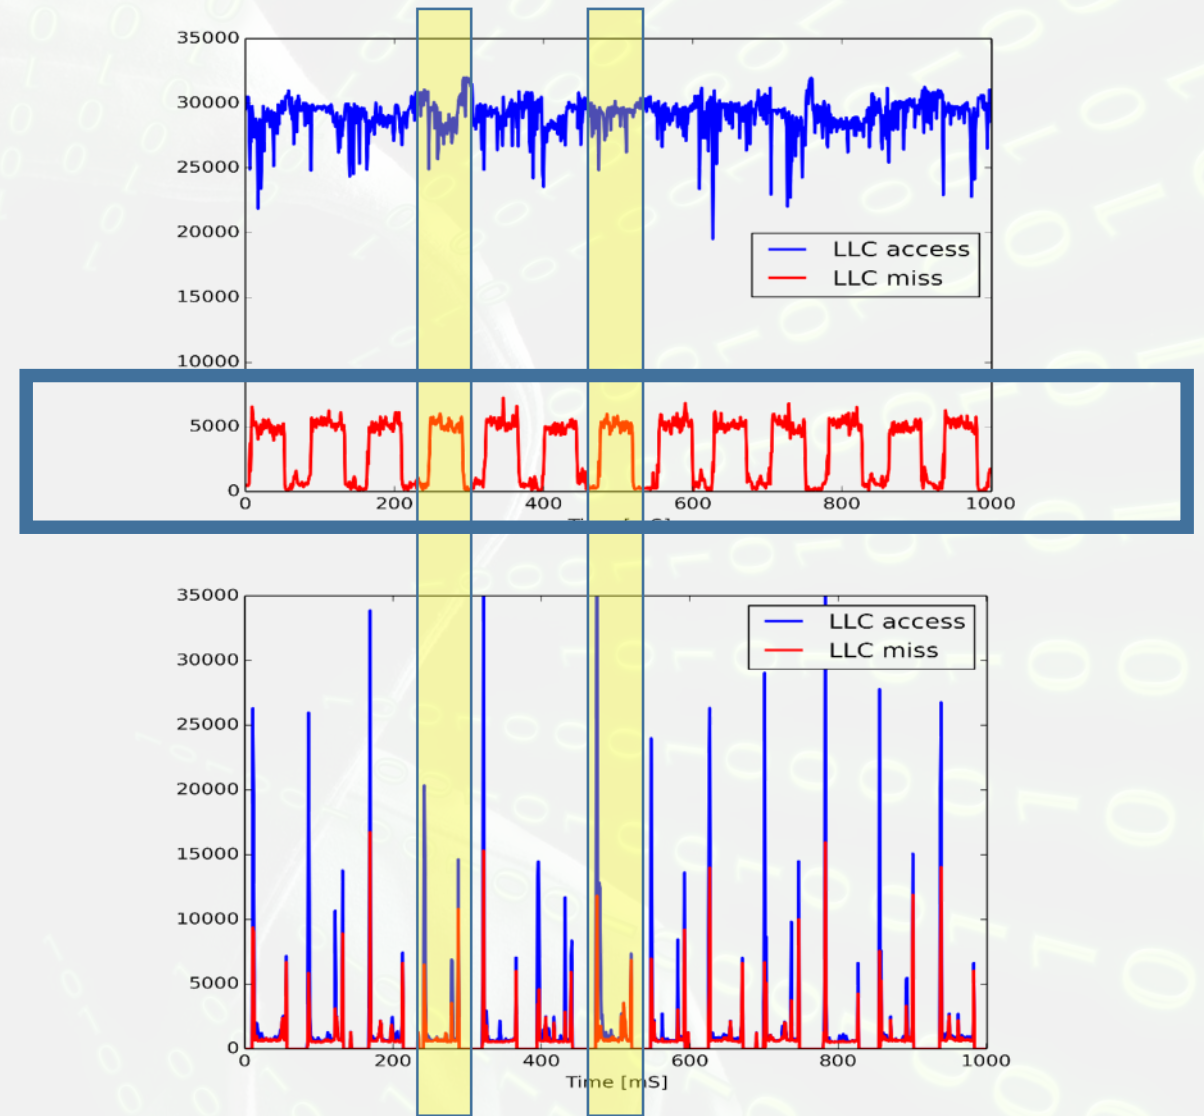
Not so nice...





# Monitoring: Active Attacker

- Can this be identified by focusing on target alone?
  - If so, monitoring could be provided as-a-service...



# Monitoring: Target-side Only (During Attack)



Sample taken every 100 $\mu$ s, MA window of 50 samples. Detection in 30ms < single RSA decryption  $\sim$ 70ms.



# And Now for Real...

- Cache side-channel attack used to extract information
- New attacks on many chip architectures (x86, ARM,...)
- Heavily based on mechanisms like the one we saw:
  - Prime&Probe
    - Separate cores, only LLC, very noisy
  - Flush&Reload
    - Attacker/target share memory, e.g., run on same core
    - Significantly easier than Prime&Probe
- Exploit various mechanisms
  - Shared memory / LLC
  - Out of order execution
    - Specifically, speculative execution
- Buzzwords: Meltdown, Spectre, ...
  - And more

## Meltdown and Spectre

Vulnerabilities in modern computers leak passwords and sensitive data.

Meltdown and Spectre exploit critical vulnerabilities in modern processors. These hardware vulnerabilities allow programs to steal data which is currently processed on the computer. While programs are typically not permitted to read data from other programs, a malicious program can exploit Meltdown and Spectre to get hold of secrets stored in the memory of other running programs. This might include your passwords stored in a password manager or browser, your personal photos, emails, instant messages and even business critical documents.

Meltdown and Spectre work on personal computers, mobile devices, and in the cloud. Depending on the cloud provider's infrastructure, it might be possible to steal data from other customers.



Meltdown

Meltdown breaks the most fundamental isolation between user applications and the operating system. This attack allows a program to access the memory, and thus also the secrets, of other programs and the operating system.

If your computer has a vulnerable processor and runs an unpatched operating system, it is not safe. Sensitive information without the chance of information. This applies both to personal as well as cloud infrastructure. Luckily, there are patches against Meltdown.

Download Paper

Cite



Spectre

Spectre breaks the isolation between different applications. It allows an attacker to trick error-free programs, which follow best practices, into leaking their secrets. In fact, the safety checks of said best practices actually increase the attack surface and may make applications more susceptible to Spectre.

The security of pretty much every computer on the planet has just gotten a lot worse



5 of the biggest data breaches (1/16)

Editor's Note: Bruce Schneier is a fellow at Harvard Kennedy School of Government. The opinions expressed in the commentary are his.

(CNN) — The security of pretty much every computer on the planet has just gotten a lot worse, and the only real solution — which of course, is not a solution — is to throw them all away and buy new ones that may be available in a few years.

On Wednesday, researchers announced a series of major security vulnerabilities in the microprocessors at the heart of the world's computers for the past 35 to 20 years. They've been





# (Partial) Bibliography

- Ristenpart et al., “Hey you get off my cloud: Exploring information leakage in third-party compute clouds”, in CCS, pp. 199-212, 2009
- Yarom and Falkner, “FLUSH+RELOAD: A High Resolution, Low Noise, L3 Cache Side-Channel Attack”. USENIX Security, pp. 719-732, 2014
- Liu et al, “Last-Level Cache Side-Channel Attacks are Practical”, IEEE S&P, pp. 605-622, 2015
- Martins et al., “A Survey on Fully Homomorphic Encryption: An Engineering Perspective”, ACM Comp. Surv. 50(6): 83:1-83:33, 2017
- Kotcher et al., “Spectre Attacks: Exploiting Speculative Execution”, ArXiv e-prints, 2018
- Lipp et al., “Meltdown”, ArXiv e-prints, 2018
- Schwarz et al., “ZombieLoad: Cross-Privilege-Boundary Data Sampling”, ArXiv e-prints, 2019
- Joint work with Niv Gilboa, Ben Amos, Arbel Levy (and others)